

MPI Exercises I

- 1 Write a script that will have each processor randomly take a sample of size 1 of TRUE and FALSE. Have each processor print its result.
- 2 Modify the script in Exercise 1 above to determine if any processors sampled TRUE. Do the same to determine if all processors sampled TRUE. In each case, print the result. Compare to the functions `comm.all()` and `comm.any()`.
- 3 Generate 50,000,000 (total) random normal values in parallel on 2, 4, and 8 processors. Time each run.
- 4
- 5 Distribute the matrix `x <- matrix(1:24, nrow=12)` in GBD format across 4 processors and call it `x.spm`.

MPI Exercises II

- 1 Add `x.spm` to itself.
- 2 Compute the mean of `x.spm`.
- 3 Compute the column means of `x.spm`.

6

7

DMAT Exercises I

- ① Subsetting, selection, and filtering are basic matrix operations featured in R. The following may look silly, but it is useful for data processing. Let `x.dmat <- ddmatrix(1:30, 10, 3)`. Do the following:

- `y.dmat <- x.dmat[c(1, 5, 4, 3),]`
`y.dmat <- x.dmat[c(10:3, 5, 5),]`
`y.dmat <- x.dmat[1:5, 3:1]`
- `y.dmat <- x.dmat[x.dmat[, 2] > 13,]`
`y.dmat <- x.dmat[x.dmat[, 2] > x.dmat[, 3],]`
`y.dmat <- x.dmat[, x.dmat[2,] > x.dmat[3,]]`
`y.dmat <- x.dmat[c(1, 3, 5), x.dmat[, 2] > x.dmat[, 3]]`

DMAT Exercises II

- ② The method `crossprod()` is an optimized form of the crossproduct computation `t(x.dmat) %*% x.dmat`. For this exercise, let `x.dmat <- ddmatrix(1:30, nrow=10, ncol=3)`.
 - ① Verify that these computations really do produce the same results.
 - ② Time each operation. Which is faster?
- ③ The `prcomp()` method returns rotations for all components. Computationally verify by example that these rotations are orthogonal, i.e., that their crossproduct is the identity matrix.
- ④