

pbdr: Bringing R Analytics to Large Distributed Architectures

George Ostrouchov

Oak Ridge National Laboratory, Oak Ridge, TN
and

University of Tennessee, Knoxville, TN

Swiss Supercomputing Center

Lugano, June 20, 2014



The pbdr Core Team

Wei-Chen Chen¹

George Ostrouchov^{2,3}

Pragneshkumar Patel³

Drew Schmidt³



Support

This work used resources of [National Institute for Computational Sciences](#) at the University of Tennessee, Knoxville, which is supported by the Office of Cyberinfrastructure of the U.S. National Science Foundation under Award No. ARRA-NSF-OCI-0906324 for NICS-RDAV center. This work also used resources of the [Oak Ridge Leadership Computing Facility](#) at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.

¹Department of Ecology and Evolutionary Biology
University of Tennessee, Knoxville TN, USA

²Computer Science and Mathematics Division
Oak Ridge National Laboratory, Oak Ridge TN, USA

³ Joint Institute for Computational Sciences
University of Tennessee, Knoxville TN, USA

Contents

1 Introduction to R

2 pbdR

3 Benchmarks

4 Applications

5 Challenges and Future

What is R?

What is R?

- *Lingua franca* for data analytics and statistical computing
 - Part programming language, part data analysis package, syntax designed for data
 - Dialect of S (Bell Labs, 1976)
 - Licensed under GPL in 1995
 - 2+ million users, dominates new work in Statistics
 - Switzerland: highest per-capita use of R!



<http://blog.revolutionanalytics.com/2014/04/a-world-map-of-r-user-activity.html>

1

Who uses R?



MERCK

KICKSTARTER

Bank of America



The
New York
Times



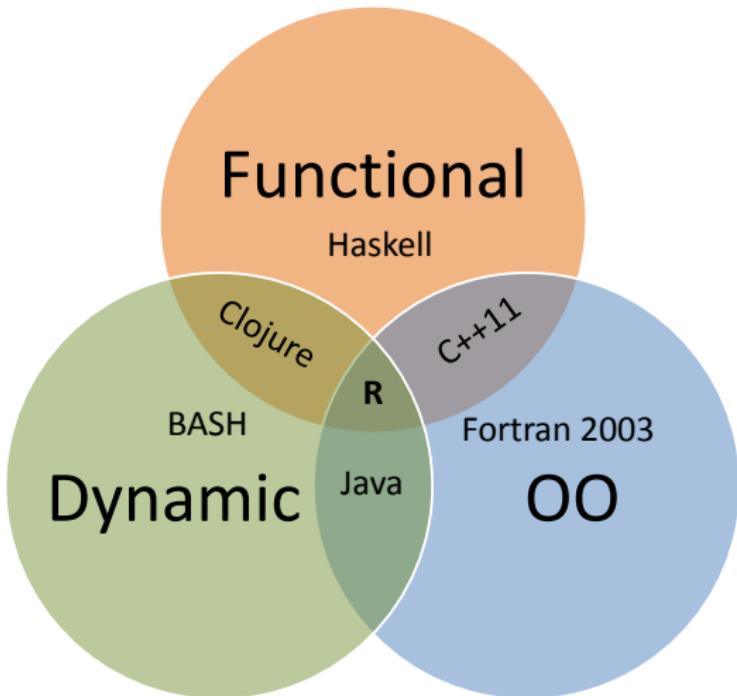
ORBITZ



Who integrates R?

- SAS
- Omniscope
- XCEL
- Tableau
- Python
- ...
- IBM-SPSS
- Java
- VTK (Paraview, VisIt)
- ORACLE
- alteryx
- KNIME

Language Paradigms



Dynamic Data Types

- Storage: logical, int, double, double complex, character
- Structures: vector, matrix, array, list, dataframe
- Caveats: (Logical) TRUE, FALSE, NA

High Level Syntax

```
1 x <- matrix(rnorm(30), nrow=10)
2 x <- x[-1, 2:5]
3 xtx <- t(x) %*% x
4 xtx <- crossprod(x)
5 ans <- prcomp(log(abs(x) + 1))
6 ans <- glm(y ~ a + b + ns(u,12), binomial(), mydata)
7 plot(ans)
8 print(ans)
```

CRAN Task View (5,000+ packages) <http://cran.rstudio.com/web/views/>

Bayesian	Bayesian Inference
ChemPhys	Chemometrics and Computational Physics
ClinicalTrials	Clinical Trial Design, Monitoring, and Analysis
Cluster	Cluster Analysis & Finite Mixture Models
DifferentialEquations	Differential Equations
Distributions	Probability Distributions
Econometrics	Computational Econometrics
Environmetrics	Analysis of Ecological and Environmental Data
ExperimentalDesign	Design of Experiments (DoE) & Analysis of Experimental Data
Finance	Empirical Finance
Genetics	Statistical Genetics
Graphics	Graphic Displays & Dynamic Graphics & Graphic Devices & Visualization
HighPerformanceComputing	High-Performance and Parallel Computing with R
MachineLearning	Machine Learning & Statistical Learning
MedicalImaging	Medical Image Analysis
MetaAnalysis	Meta-Analysis
Multivariate	Multivariate Statistics
NaturalLanguageProcessing	Natural Language Processing
NumericalMathematics	Numerical Mathematics
OfficialStatistics	Official Statistics & Survey Methodology
Optimization	Optimization and Mathematical Programming
Pharmacokinetics	Analysis of Pharmacokinetic Data
Phylogenetics	Phylogenetics, Especially Comparative Methods
Psychometrics	Psychometric Models and Methods
ReproducibleResearch	Reproducible Research
Robust	Robust Statistical Methods
SocialSciences	Statistics for the Social Sciences
Spatial	Analysis of Spatial Data
SpatioTemporal	Handling and Analyzing Spatio-Temporal Data
Survival	Survival Analysis
TimeSeries	Time Series Analysis
WebTechnologies	Web Technologies and Services
gR	gRaphical Models in R

Other R package repositories



- Tools for the analysis and comprehension of high-throughput genomic data
- 800+ packages <http://www.bioconductor.org>
- Rmetrics
 - 30+ packages for financial computing
<https://rmetrics.org> (ETH)
- ...

Why take R to HPC?

- ① People love it
 - ② Highly extensible
 - ③ Unmatched diversity of analytical methods
 - ④ Thorough treatment of uncertainty throughout
 - ⑤ Gold standard for graphics (traditional, grid, and ggplot2)
 - ⑥ HPC community needs a high-level language for data analysis

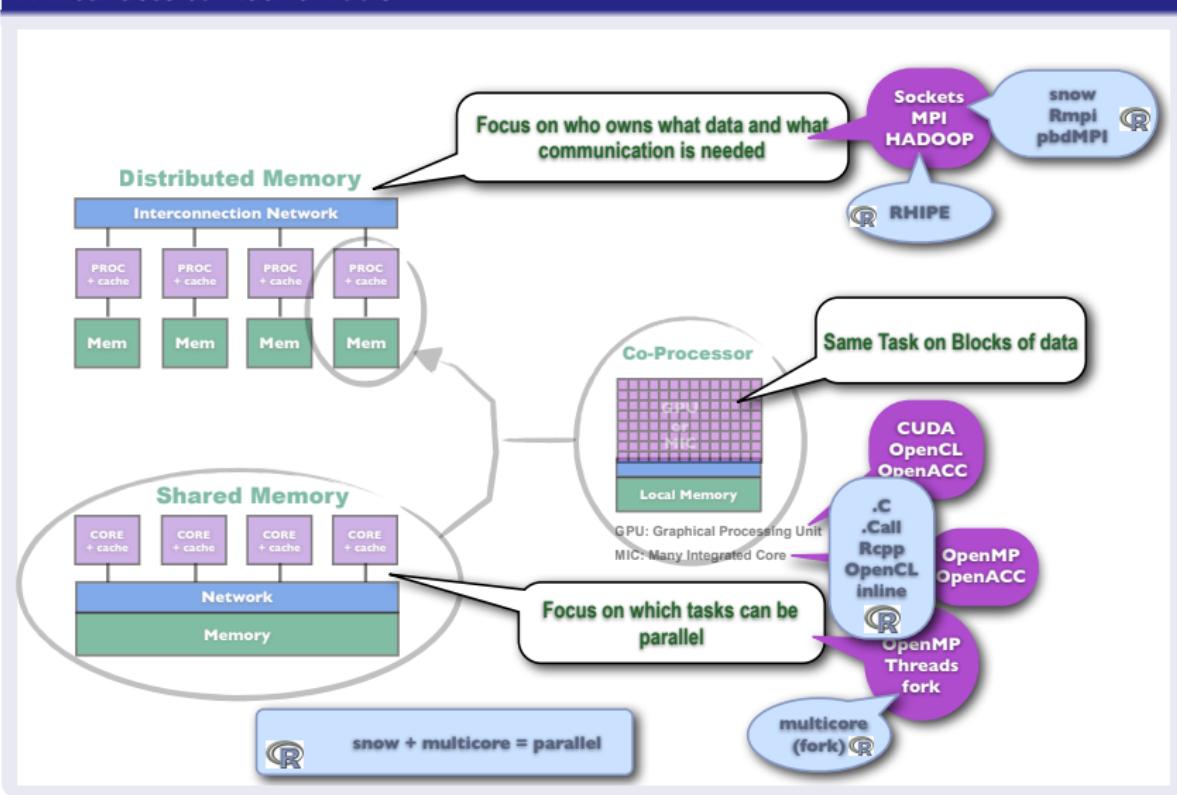
Data analysis is interactive!

- Data reduction to knowledge
- S (and R) interactive “answer” to batch data analysis
- Efficient use of expensive people
- Iterative process with same data
 - Diagnostics of fit
 - Quantification of uncertainty
 - Interpretation of results

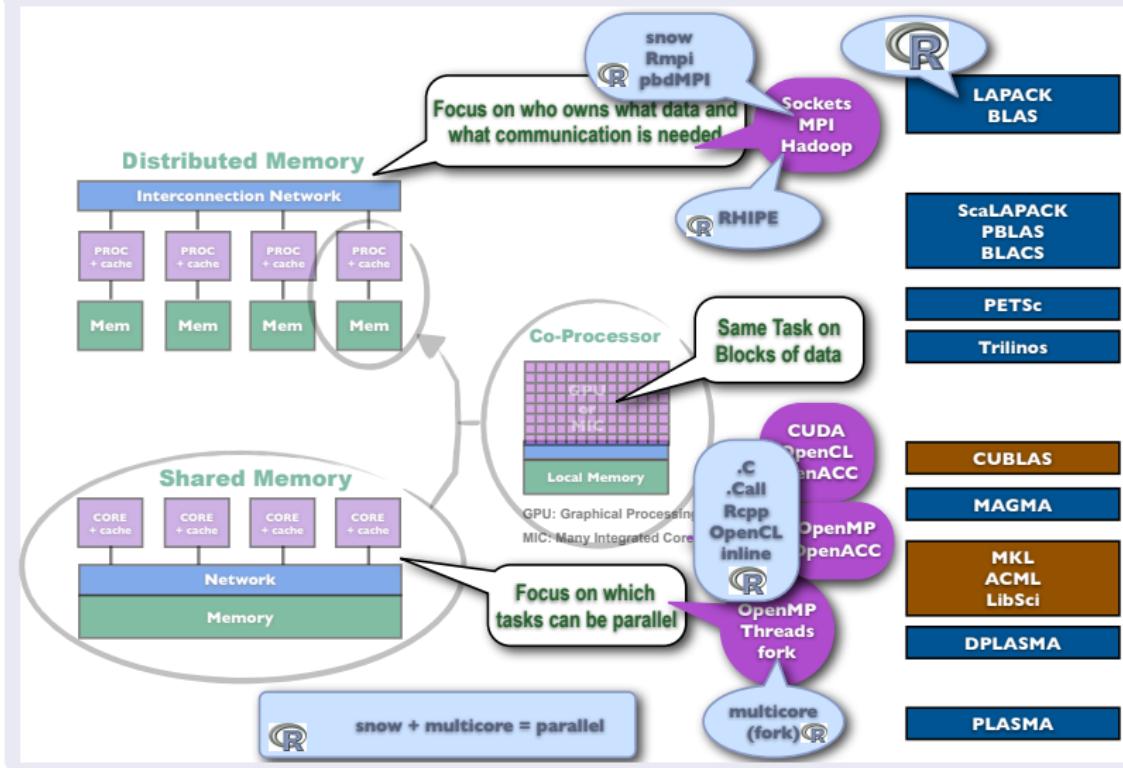
HPC is batch!

- Traditionally data generation
- Efficient use of expensive platforms
- Long time to assimilate generated data

R Interfaces to Native Tools

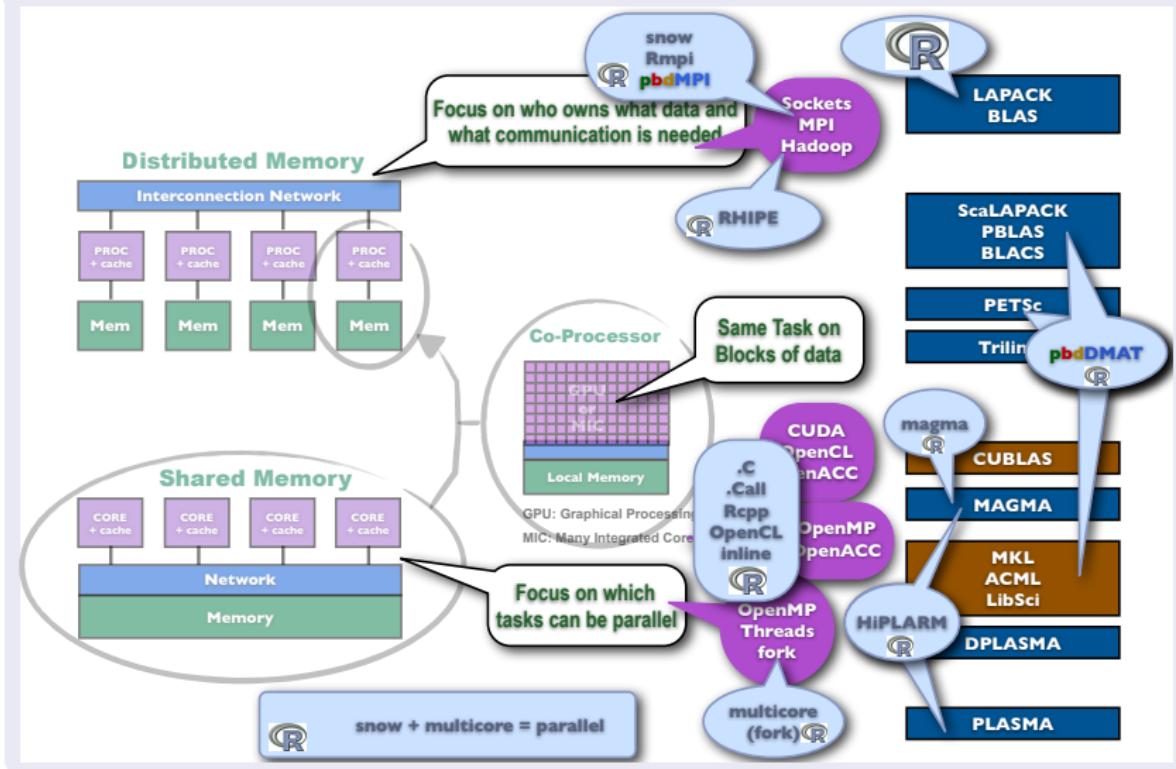


Scalable Math Libraries from 30+ Years of Research



R and HPC

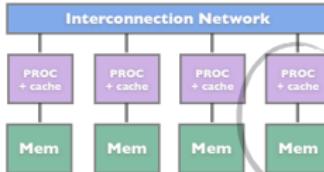
R Interfaces to Scalable Math Libraries



R Interfaces to Scalable Math Libraries

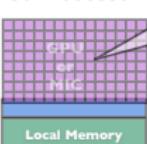


Distributed Memory



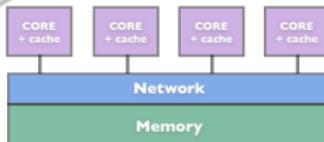
Focus on who owns what data and what communication is needed

Co-Processor



Same Task on Blocks of data

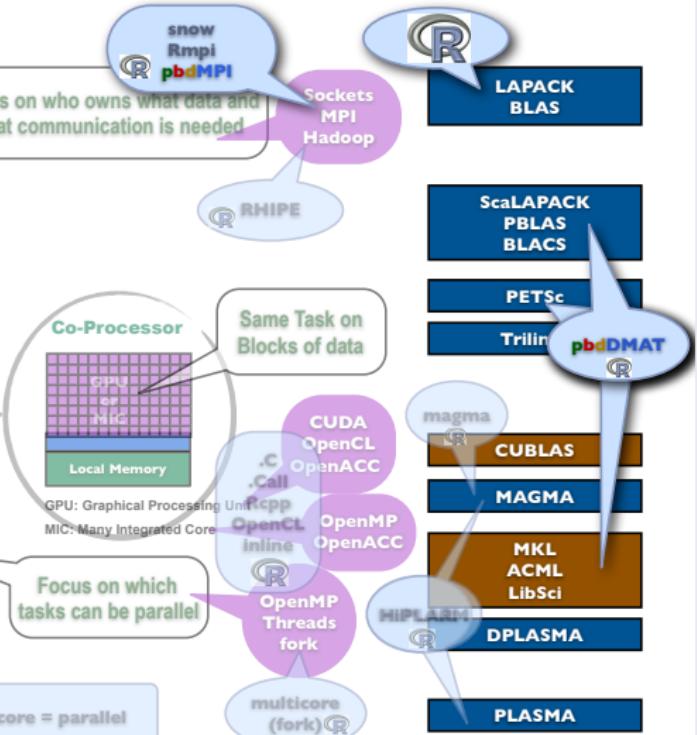
Shared Memory



Focus on which tasks can be parallel



snow + multicore = parallel



Contents

1 Introduction to R

2 pbdR

3 Benchmarks

4 Applications

5 Challenges and Future

Programming with Big Data in R (**pbDR**)

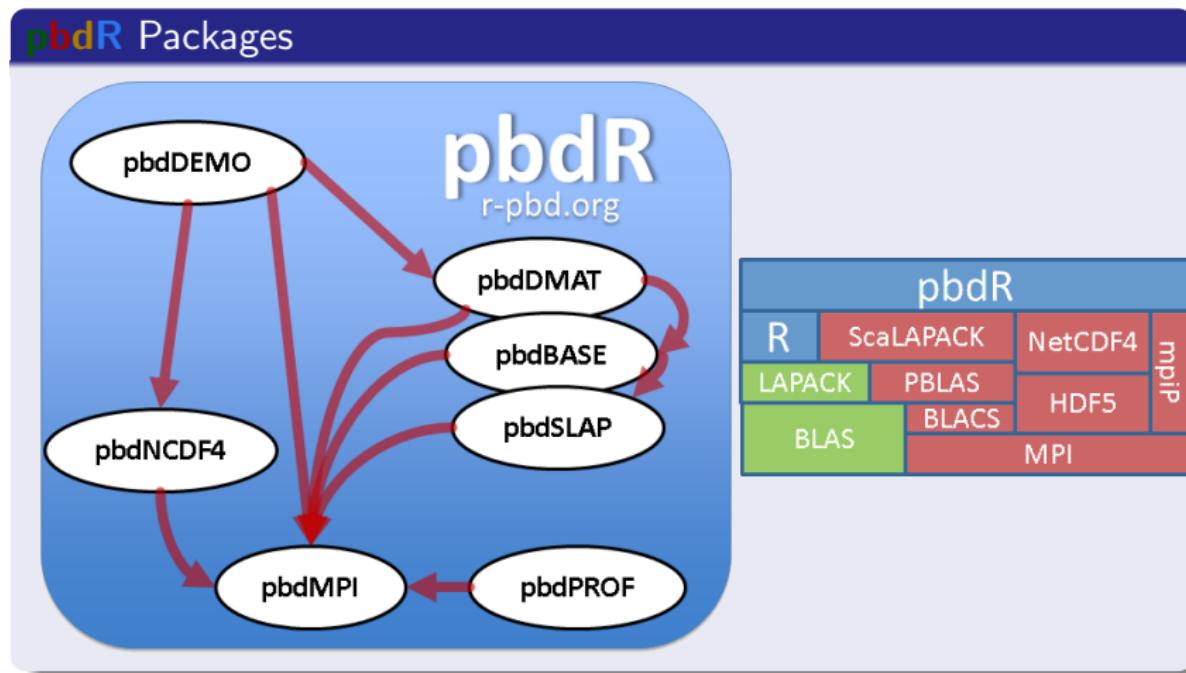
Productivity, Portability, Performance



- Free^a R packages.
- Bridging high-performance C with high-productivity of R
- Distributed data details implicitly managed.
- Methods have syntax *identical* to R.

^aMPL, BSD, and GPL licensed

The pbdr Project



pbdmPI vs Rmpi: API

Reduction Operations

Rmpi

```
1 # int
2 mpi.allreduce(x, type=1)
3 # double
4 mpi.allreduce(x, type=2)
```

pbdmPI

```
1 allreduce(x)
```

Types in R

```
1 > is.integer(1)
2 [1] FALSE
3 > is.integer(2)
4 [1] FALSE
5 > is.integer(1:2)
6 [1] TRUE
```

MPI Operations: The Gang's All Here

- **Communicator wrangling:** `init()`, `finalize()`
- **Rank query:** `comm.rank()`, `comm.size()`
- **Reduction:** `reduce(x, op='sum')`, `allreduce(x)`
- **Gather:** `gather(x)`, `allgather(x)`
- **Broadcast:** `bcast(x)`
- **Barrier:** `barrier()`
- **Send/Receive:** `send(x)`, `recv()`

MPI Operations for R Users

- **Printing:** comm.print(x), comm.cat(x)
- **RNG Seeds:** comm.set.seed(diff=T)
- **Task Subsetting:** get.jid(n)
- ***ply:**
`pbdrApply(X, MARGIN, FUN, ...)`
`pbdrLapply(X, FUN, ...)`
`pbdrSapply(X, FUN, ...)`

Rank Query Example

```
1 library(pbdrMPI, quiet = TRUE)
2 init()
3
4 my.rank <- comm.rank()
5 comm.print(my.rank, all.rank=TRUE)
6
7 finalize()
```

Execute this script via:

```
1 mpirun -np 2 Rscript 1_rank.r
```

Sample Output:

```
1 COMM.RANK = 0
2 [1] 0
3 COMM.RANK = 1
4 [1] 1
```

2×3 block-cyclic grid on 6 processors: Global view

$$x = \begin{bmatrix} x_{11} & x_{12} & x_{13} & x_{14} & x_{15} & x_{16} & x_{17} & x_{18} & x_{19} \\ x_{21} & x_{22} & x_{23} & x_{24} & x_{25} & x_{26} & x_{27} & x_{28} & x_{29} \\ \hline x_{31} & x_{32} & x_{33} & x_{34} & x_{35} & x_{36} & x_{37} & x_{38} & x_{39} \\ x_{41} & x_{42} & x_{43} & x_{44} & x_{45} & x_{46} & x_{47} & x_{48} & x_{49} \\ \hline x_{51} & x_{52} & x_{53} & x_{54} & x_{55} & x_{56} & x_{57} & x_{58} & x_{59} \\ x_{61} & x_{62} & x_{63} & x_{64} & x_{65} & x_{66} & x_{67} & x_{68} & x_{69} \\ \hline x_{71} & x_{72} & x_{73} & x_{74} & x_{75} & x_{76} & x_{77} & x_{78} & x_{79} \\ x_{81} & x_{82} & x_{83} & x_{84} & x_{85} & x_{86} & x_{87} & x_{88} & x_{89} \\ \hline x_{91} & x_{92} & x_{93} & x_{94} & x_{95} & x_{96} & x_{97} & x_{98} & x_{99} \end{bmatrix}_{9 \times 9}$$

$$\text{Processor grid} = \left| \begin{array}{ccc} 0 & 1 & 2 \\ 3 & 4 & 5 \end{array} \right| = \left| \begin{array}{ccc} (0,0) & (0,1) & (0,2) \\ (1,0) & (1,1) & (1,2) \end{array} \right|$$

2×3 block-cyclic grid on 6 processors: Local view

X_{11}	X_{12}	X_{17}	X_{18}	5×4	X_{13}	X_{14}	X_{19}	5×3	X_{15}	X_{16}	5×2
X_{21}	X_{22}	X_{27}	X_{28}		X_{23}	X_{24}	X_{29}		X_{25}	X_{26}	
X_{51}	X_{52}	X_{57}	X_{58}		X_{53}	X_{54}	X_{59}		X_{55}	X_{56}	
X_{61}	X_{62}	X_{67}	X_{68}		X_{63}	X_{64}	X_{69}		X_{65}	X_{66}	
X_{91}	X_{92}	X_{97}	X_{98}		X_{93}	X_{94}	X_{99}		X_{95}	X_{96}	
X_{31}	X_{32}	X_{37}	X_{38}	4×4	X_{33}	X_{34}	X_{39}	4×3	X_{35}	X_{36}	4×2
X_{41}	X_{42}	X_{47}	X_{48}		X_{43}	X_{44}	X_{49}		X_{45}	X_{46}	
X_{71}	X_{72}	X_{77}	X_{78}		X_{73}	X_{74}	X_{79}		X_{75}	X_{76}	
X_{81}	X_{82}	X_{87}	X_{88}		X_{83}	X_{84}	X_{89}		X_{85}	X_{86}	

Processor grid =
$$\begin{vmatrix} 0 & 1 & 2 \\ 3 & 4 & 5 \end{vmatrix} = \begin{vmatrix} (0,0) & (0,1) & (0,2) \\ (1,0) & (1,1) & (1,2) \end{vmatrix}$$

pbDR Example Syntax

```
1 x <- x[-1, 2:5]
2 x <- log(abs(x) + 1)
3 xtx <- t(x) %*% x
4 ans <- svd(solve(xtx))
```

The above runs on 1 core with R or 10,000 cores with pbDR

Profiling with pbdPROF

1. Rebuild **pbdR** packages

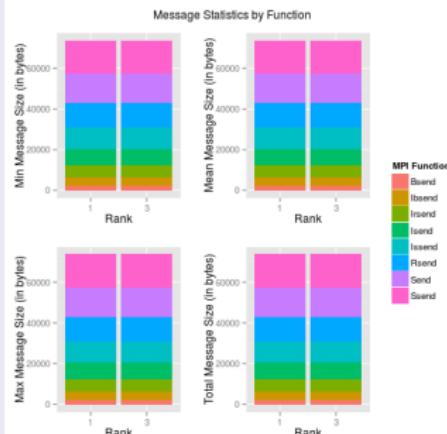
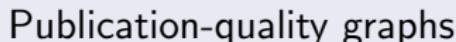
```
R CMD INSTALL  
    pbdMPI_0.2-1.tar.gz \  
    --configure-args= \  
    "--enable-pbdPROF"
```

2. Run code

```
mpirun -np 64 Rscript  
my_script.R
```

3. Analyze results

```
1 library(pbdPROF)
2 prof <- read.prof(
3     "profiler_output.mpiP")
4 plot(prof)
```



pbdPAPI

Example

Contents

1 Introduction to R

2 pbdR

3 Benchmarks

4 Applications

5 Challenges and Future

Non-Optimal Choices Throughout

- ① Only free software used (no MKL, ACML, etc.)
- ② 1 core = 1 MPI process
- ③ No tuning for data distribution, just the defaults

Benchmark Data

- ① Measure wallclock time for covariance and linear regression
- ② Random normal $N(100, 10000)$
- ③ Local problem size fixed at $\approx 43.4MiB$
- ④ “weak scaling” = global problem grows with core count
- ⑤ Three sets: 500, 1000, and 2000 columns
- ⑥ Several runs at different core counts within each set



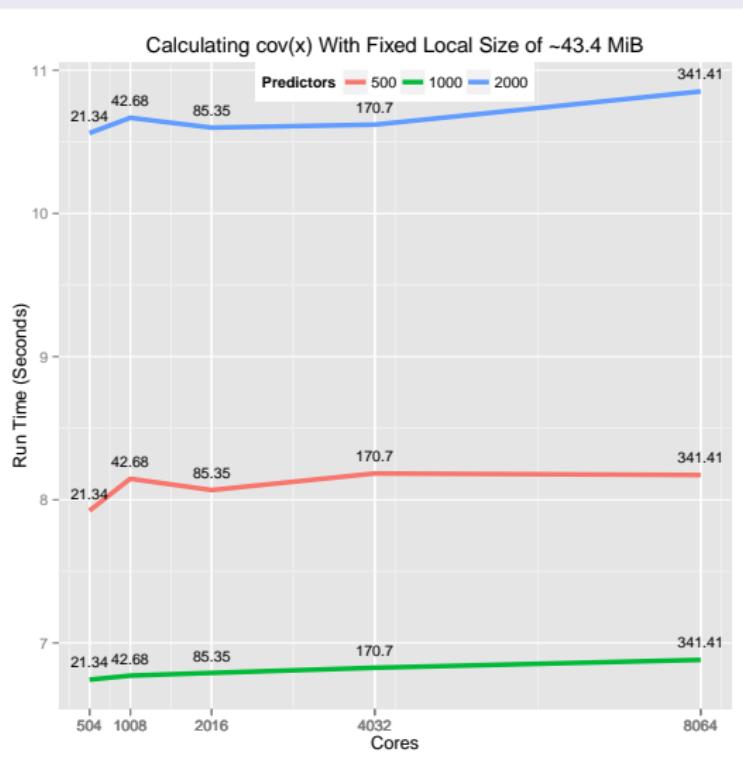
Covariance Code

$$\text{cov}(x_{n \times p}) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu_x)(x_i - \mu_x)^T$$

```
1 x <- ddmatrix("rnorm", nrow=n, ncol=p, mean=mean, sd=sd)  
2  
3 cov.x <- cov(x)
```

Benchmarks

cov()



Benchmarks

Linear Model Code

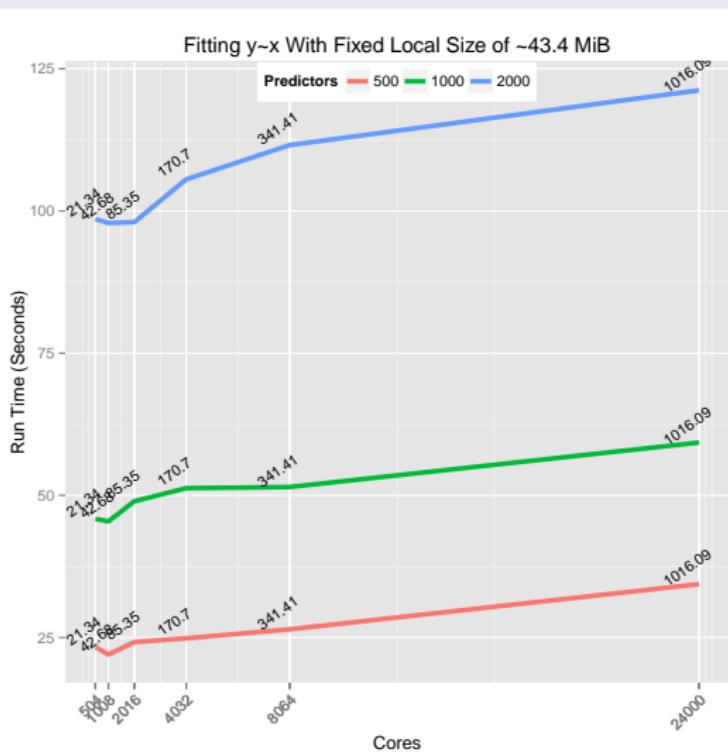
Find β such that

$$y = X\beta + \epsilon$$

When X is full rank,

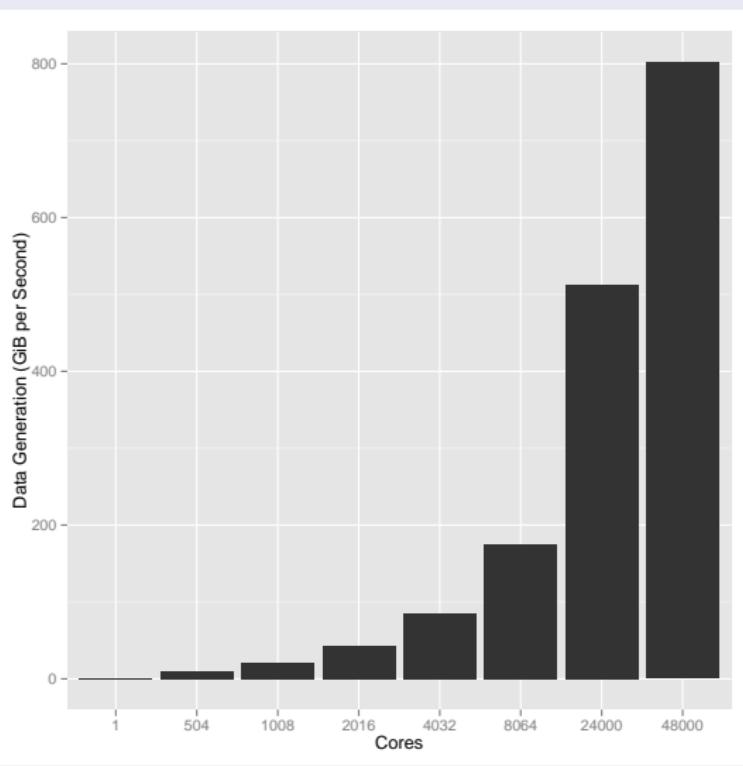
$$\hat{\beta} = (X^T X)^{-1} X^T y$$

```
1 x <- ddmatrix("rnorm", nrow=n, ncol=p, mean=100,  
                 sd=10000)  
2 beta_true <- ddmatrix("runif", nrow=p, ncol=1)  
3  
4 y <- x %*% beta_true  
5  
6 beta_est <- lm.fit(x, y)$coefficients
```



Benchmarks

Data Generation



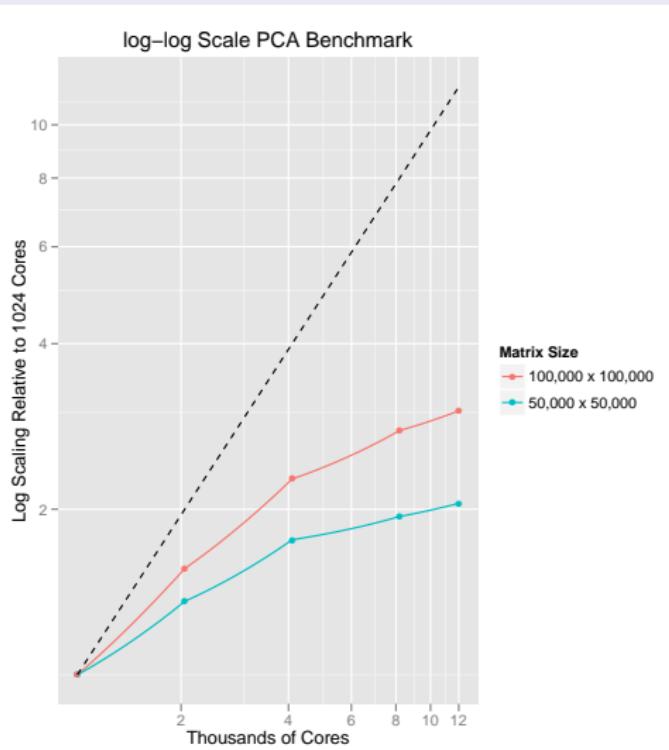
PCA Benchmark Data

- ① Measure wallclock time for principal components analysis
- ② Random normal $N(100, 10000)$
- ③ Global problem size fixed
- ④ “strong scaling” = local problem decreases with core count
- ⑤ Two sets: $50,000 \times 50,000$ and $100,000 \times 100,000$
- ⑥ Runs out of local work as core count increases



Benchmarks

prcomp() Strong Scaling



Contents

1 Introduction to R

2 pbdR

3 Benchmarks

4 Applications

5 Challenges and Future

Productivity Example

Randomized SVD¹

PROTOTYPE FOR RANDOMIZED SVD

Given an $m \times n$ matrix A , a target number k of singular vectors, and an exponent q (say, $q = 1$ or $q = 2$), this procedure computes an approximate rank- $2k$ factorization $U\Sigma V^*$, where U and V are orthonormal, and Σ is nonnegative and diagonal.

Stage A:

- 1 Generate an $n \times 2k$ Gaussian test matrix Ω .
- 2 Form $Y = (AA^*)^q A\Omega$ by multiplying alternately with A and A^* .
- 3 Construct a matrix Q whose columns form an orthonormal basis for the range of Y .

Stage B:

- 4 Form $B = Q^*A$.
- 5 Compute an SVD of the small matrix: $B = \tilde{U}\Sigma V^*$.
- 6 Set $U = Q\tilde{U}$.

Note: The computation of Y in step 2 is vulnerable to round-off errors. When high accuracy is required, we must incorporate an orthonormalization step between each application of A and A^* ; see Algorithm 4.4.

ALGORITHM 4.4: RANDOMIZED SUBSPACE ITERATION

Given an $m \times n$ matrix A and integers ℓ and q , this algorithm computes an $m \times \ell$ orthonormal matrix Q whose range approximates the range of A .

- 1 Draw an $n \times \ell$ standard Gaussian matrix Ω .
- 2 Form $Y_0 = A\Omega$ and compute its QR factorization $Y_0 = Q_0R_0$.
- 3 **for** $j = 1, 2, \dots, q$
- 4 Form $\tilde{Y}_j = A^*Q_{j-1}$ and compute its QR factorization $\tilde{Y}_j = \tilde{Q}_j\tilde{R}_j$.
- 5 Form $Y_j = A\tilde{Q}_j$ and compute its QR factorization $Y_j = Q_jR_j$.
- 6 **end**
- 7 $Q = Q_q$.

Serial R

```

1 randSVD <- function(A, k, q=3)
2 {
3   ## Stage A
4   Omega <- matrix(rnorm(n*2*k),
5                     nrow=n, ncol=2*k)
6   Y <- A %*% Omega
7   Q <- qr.Q(qr(Y))
8   At <- t(A)
9   for(i in 1:q)
10  {
11    Y <- At %*% Q
12    Q <- qr.Q(qr(Y))
13    Y <- A %*% Q
14    Q <- qr.Q(qr(Y))
15  }
16
17  ## Stage B
18  B <- t(Q) %*% A
19  U <- La.svd(B)$u
20  U <- Q %*% U
21  U[, 1:k]
22 }
```

¹Halko N, Martinsson P-G and Tropp J A 2011 Finding structure with randomness: probabilistic algorithms for constructing approximate matrix decompositions *SIAM Rev.* **53** 217–88

Productivity Example

Randomized SVD

Serial R

```

1 randSVD <- function(A, k, q=3)
2 {
3   ## Stage A
4   Omega <- matrix(rnorm(n*2*k),
5                     nrow=n, ncol=2*k)
6   Y <- A %*% Omega
7   Q <- qr.Q(qr(Y))
8   At <- t(A)
9   for(i in 1:q)
10  {
11    Y <- At %*% Q
12    Q <- qr.Q(qr(Y))
13    Y <- A %*% Q
14    Q <- qr.Q(qr(Y))
15  }
16
17   ## Stage B
18   B <- t(Q) %*% A
19   U <- La.svd(B)$u
20   U <- Q %*% U
21   U[, 1:k]
22 }
```

Parallel pbdr

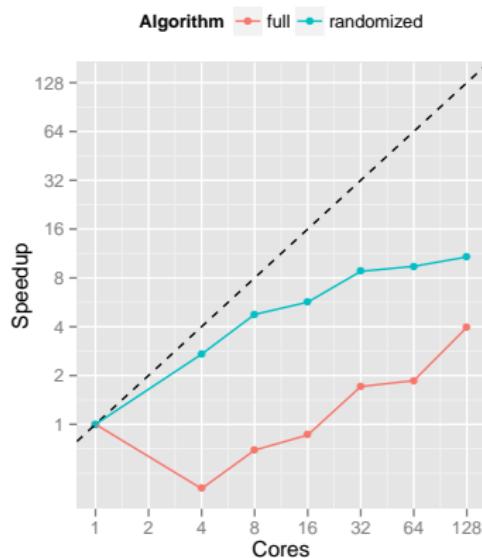
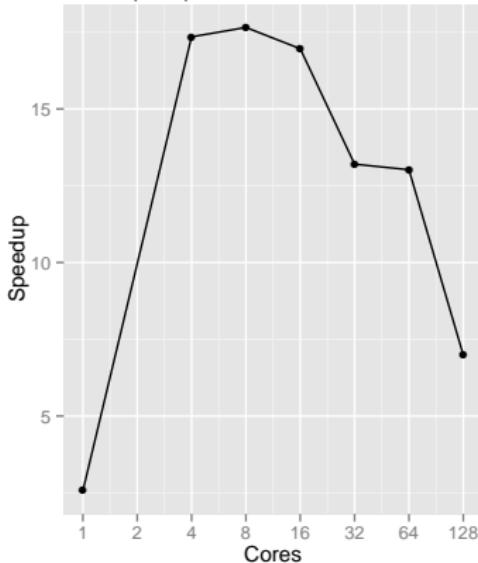
```

1 randSVD <- function(A, k, q=3)
2 {
3   ## Stage A
4   Omega <- ddmatrix("rnorm",
5                     nrow=n, ncol=2*k)
6   Y <- A %*% Omega
7   Q <- qr.Q(qr(Y))
8   At <- t(A)
9   for(i in 1:q)
10  {
11    Y <- At %*% Q
12    Q <- qr.Q(qr(Y))
13    Y <- A %*% Q
14    Q <- qr.Q(qr(Y))
15  }
16
17   ## Stage B
18   B <- t(Q) %*% A
19   U <- La.svd(B)$u
20   U <- Q %*% U
21   U[, 1:k]
22 }
```

Productivity Example

Randomized SVD on ≈ 765 MiB

30 Singular Vectors from a 100,000 by 1,000 Matrix

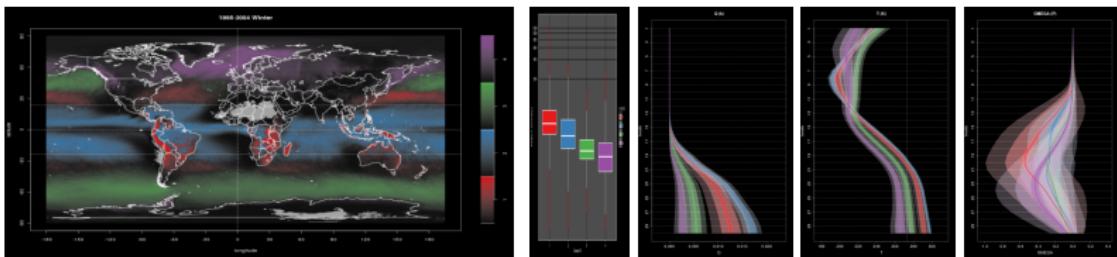
30 Singular Vectors from a 100,000 by 1,000 Matrix
Speedup of Randomized vs. Full SVD

Science Applications

Exploration of Climate Data (2.9 TB) with pbdR

- 2013 INCITE “Attributing Changes in the Risk of Extreme Weather and Climate” Michael Wehner, PI
 - High resolution atmospheric model, CAM 5.1 (~25 km resolution)
 - Resolves processes responsible for extreme precipitation and storms
- R and pbdR used for scalable analytics and graphics
 - Advanced clustering capability (EM algorithm for Gaussian mixture models)

Behavior of extreme precipitation across atmospheric layers of moisture, temperature, and wind

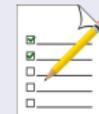


W.-C. Chen, G. Ostrochov, D. Pugmire, Prabhat, M. Wehner. A Parallel EM Algorithm for Model-Based Clustering Applied to the Exploration of Large Spatio-Temporal Data. *Technometrics* (online, in print: Fall 2013).

Science Applications

Simulation in the Social Sciences with pbdr

- Survey responses can be intentionally misleading
- This results in bias and poor policy decisions
- Simulation examines bias



*“... [Serial] simulations would require nearly a year to complete.
... Using [pbdr] and NICS resources, the simulation completed
in just 15 hours.”*

— Joseph Robinson, Educational Psychology, University of Illinois at Urbana-Champaign

Engaging with the R Community

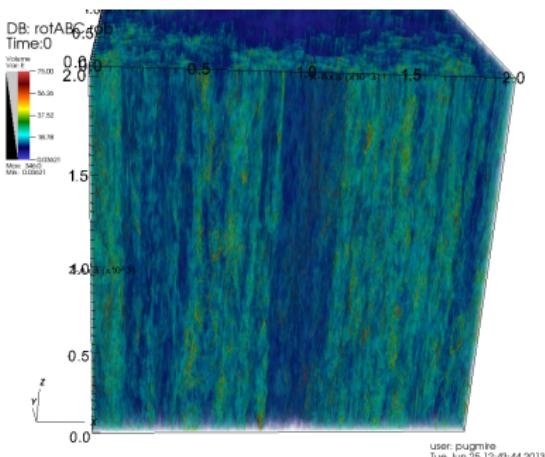
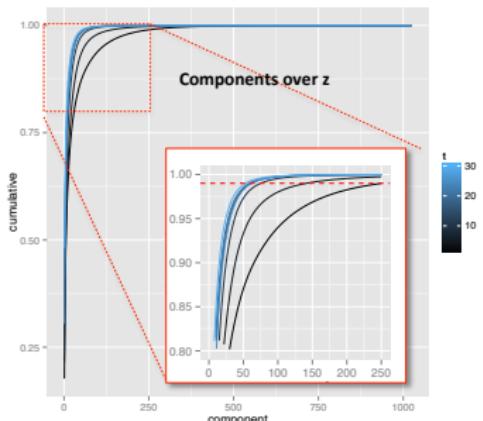
*“It’s exactly what I wanted ever since I started learning about HPC
with C and MPI a few years back!”*

— Andrew Raim, Department of Mathematics and Statistics,
University of Maryland, Baltimore County

Science Applications

Coherent Structures in Turbulent Flow with **pbdR**

- Data: Turbulence cube: $2048^3 \approx 8.6$ Billion elements ≈ 70 GiB per time step
 - Principal components analysis across z
 - Subset: $512 \times 512 \times 1,024$
 - I/O + subset + PCA + plots ≈ 60 seconds per time step with 128 cores



- **Coherence in z increases over time:**
 - *Time 1: 25% components capture 99% variability*
 - *Time 30: 5% components capture 99% variability*

2013 INCITE "Parameter Studies of Boussinesq Flows" Susan Kurien, PI, Duane Rosenberg, Co-PI, et al.

Contents

1 Introduction to R

2 pbdR

3 Benchmarks

4 Applications

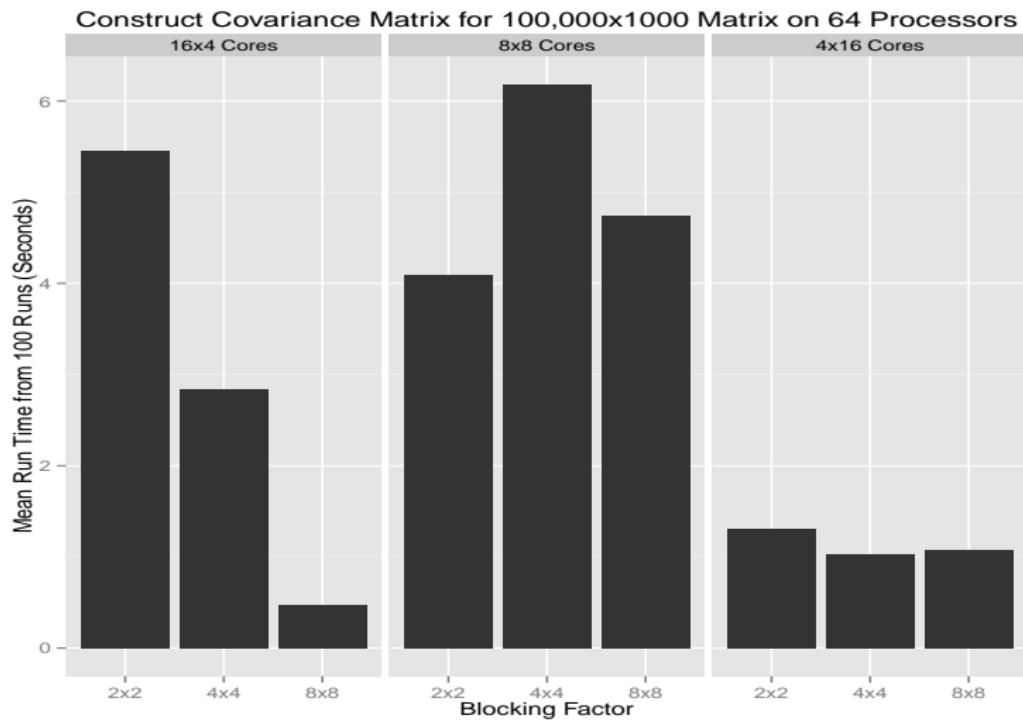
5 Challenges and Future

Challenges

Challenges

- Perceptions.
 - “R? Isn’t that slow?” – HPC people
 - “HPC? Isn’t that hard?” – R people
- Bringing R community to large platforms
- Distributed data input
- Bringing interactivity back
- Performance dependence on data layout

Covariance Revisited: Data layout parameter calibration



Future Work

Future Work

- Starting a 3 year NSF grant to work on
 - Bring back interactivity via client/server
 - Simplify parallel data input
 - Begin DPLASMA integration
- Optimization for heterogeneous architectures
- Support for more profilers
- In-situ use with simulations (QMC, AutoTune)

Future Work

Where to learn more?

- <http://r-pbd.org/>
 - **pbdDEMO** vignette
 - RBigData@gmail.com

Thanks for coming!

Questions?



pbdR
Programming with Big Data in R

<http://r-pbd.org/>