

Programming with Big Data in R

George Ostrouchov

February 17, 2014



The pbdR Core Team

Wei-Chen Chen¹
George Ostrouchov²
Pragneshkumar Patel³
Drew Schmidt³



Support

This work used resources of [National Institute for Computational Sciences](#) at the University of Tennessee, Knoxville, which is supported by the Office of Cyberinfrastructure of the U.S. National Science Foundation under Award No. ARRA-NSF-OCI-0906324 for NICS-RDAV center. This work also used resources of the [Oak Ridge Leadership Computing Facility](#) at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.

¹Department of Ecology and Evolutionary Biology
University of Tennessee, Knoxville TN, USA

²Computer Science and Mathematics Division
Oak Ridge National Laboratory, Oak Ridge TN, USA

³ National Institute for Computational Sciences
University of Tennessee, Knoxville TN, USA

About This Presentation

Speaking Serial R with a Parallel Accent

The content of this presentation is based in part on the **pbdDEMO** vignette *Speaking Serial R with a Parallel Accent*

<http://goo.gl/HZkRt>

It contains more examples, and sometimes added detail.

About This Presentation

Installation Instructions

Installation instructions for setting up a pbdr environment are available:

<http://r-pbd.org/install.html>

This includes instructions for installing R, MPI, and pbdr.

Contents

1 pbdR

Contents

1 pbdR

- The pbdR Project
- pbdR Paradigms

Message Passing Interface (MPI)

- *MPI*: Standard for managing communications (data and instructions) between different nodes/computers.
- *Implementations*: OpenMPI, MPICH2, Cray MPT, ...
- Enables parallelism (via communication) on distributed machines.
- *Communicator*: manages communications between processors.

MPI Operations (1 of 2)

- **Managing a Communicator:** Create and destroy communicators.
`init()` — initialize communicator
`finalize()` — shut down communicator(s)
- **Rank query:** determine the processor's position in the communicator.
`comm.rank()` — “who am I?”
`comm.size()` — “how many of us are there?”
- **Printing:** Printing output from various ranks.
`comm.print(x)`
`comm.cat(x)`
WARNING: only use these functions on *results*, never on yet-to-be-computed things.

Quick Example 1

Rank Query: 1_rank.r

```
1 library(pbdMPI, quiet = TRUE)
2 init()
3
4 my.rank <- comm.rank()
5 comm.print(my.rank, all.rank=TRUE)
6
7 finalize()
```

Execute this script via:

```
1 mpirun -np 2 Rscript 1_rank.r
```

Sample Output:

```
1 COMM.RANK = 0
2 [1] 0
3 COMM.RANK = 1
4 [1] 1
```

Quick Example 2

Hello World: 2_hello.r

```
1 library(pbdMPI, quiet=TRUE)
2 init()
3
4 comm.print("Hello, world")
5
6 comm.print("Hello again", all.rank=TRUE, quiet=TRUE)
7
8 finalize()
```

Execute this script via:

```
1 mpirun -np 2 Rscript 2_hello.r
```


Sample Output:

```
1 COMM.RANK = 0
2 [1] "Hello, world"
3 [1] "Hello again"
4 [1] "Hello again"
```


MPI Operations

- 1 Reduce
- 2 Gather
- 3 Broadcast
- 4 Barrier


Reductions — Combine results into single result




Gather — Many-to-one



Broadcast — One-to-many



Barrier — Synchronization



MPI Operations (2 of 2)

- **Reduction:** each processor has a number x ; add all of them up, find the largest/smallest,
`reduce(x, op='sum')` — reduce to one
`allreduce(x, op='sum')` — reduce to all
- **Gather:** each processor has a number; create a new object on some processor containing all of those numbers.
`gather(x)` — gather to one
`allgather(x)` — gather to all
- **Broadcast:** one processor has a number x that every other processor should also have.
`bcast(x)`
- **Barrier:** “computation wall”; no processor can proceed until *all* processors can proceed.
`barrier()`

Quick Example 3

Reduce and Gather: 3_gt.r

```
1 library(pbdMPI, quiet = TRUE)
2 init()
3
4 comm.set.seed(diff=TRUE)
5
6 n <- sample(1:10, size=1)
7
8 gt <- gather(n)
9 comm.print(unlist(gt))
10
11 sm <- allreduce(n, op='sum')
12 comm.print(sm, all.rank=T)
13
14 finalize()
```

Execute this script via:

```
1 mpirun -np 2 Rscript 3_gt.r
```

Sample Output:

```
1 COMM.RANK = 0
2 [1] 2 8
3 COMM.RANK = 0
4 [1] 10
5 COMM.RANK = 1
6 [1] 10
```

Quick Example 4

Broadcast: 4_bcast.r

```
1 library(pbdMPI, quiet=T)
2 init()
3
4 if (comm.rank()==0){
5   x <- matrix(1:4, nrow=2)
6 } else {
7   x <- NULL
8 }
9
10 y <- bcast(x, rank.source=0)
11
12 comm.print(y, rank=1)
13
14 finalize()
```

Execute this script via:

```
1 mpirun -np 2 Rscript 4_bcast.r
```

Sample Output:

```
1 COMM.RANK = 1
2   [,1] [,2]
3 [1,]   1   3
4 [2,]   2   4
```

MPI Package Controls

The `.SPMD.CT` object allows for setting different package options with **pbdMPI**. See the entry *SPMD Control* of the **pbdMPI** manual for information about the `.SPMD.CT` object:

<http://cran.r-project.org/web/packages/pbdMPI/pbdMPI.pdf>

Quick Example 5

Barrier: 5_barrier.r

```
1 library(pbdMPI, quiet = TRUE)
2 init()
3
4 .SPMD.CT$msg.barrier <- TRUE
5 .SPMD.CT$print.quiet <- TRUE
6
7 for (rank in 1:comm.size()-1){
8   if (comm.rank() == rank){
9     cat(paste("Hello", rank+1, "of", comm.size(), "\n"))
10  }
11  barrier()
12 }
13
14 comm.cat("\n")
15
16 comm.cat(paste("Hello", comm.rank()+1, "of",
17               comm.size(), "\n"), all.rank=TRUE)
18 finalize()
```

Execute this script via:

```
1 mpirun -np 2 Rscript 5_barrier.r
```

Sample Output:

```
1 Hello 1 of 2
2 Hello 2 of 2
```

Random Seeds

pbdMPI offers a simple interface for managing random seeds:

- `comm.set.seed(diff=TRUE)` — Independent streams via the **rlecuyer** package.
- `comm.set.seed(seed=1234, diff=FALSE)` — All processors use the same seed `seed=1234`
- `comm.set.seed(diff=FALSE)` — All processors use the same seed, determined by processor 0 (using the system clock and PID of processor 0).

Quick Example 6

Timing: 6_timer.r

```

1 library(pbdMPI, quiet=TRUE)
2 init()
3
4 comm.set.seed(diff=T)
5
6 test <- function(timed)
7 {
8   ltime <- system.time(timed)[3]
9
10  mintime <- allreduce(ltime, op='min')
11  maxtime <- allreduce(ltime, op='max')
12  meantime <- allreduce(ltime, op='sum')/comm.size()
13
14  return(data.frame(min=mintime, mean=meantime,
15                    max=maxtime))
16 }
17 times <- test(rnorm(1e6)) # ~7.6MiB of data
18 comm.print(times)
19
20 finalize()

```

Execute this script via:

```
1 mpirun -np 2 Rscript 6_timer.r
```

Sample Output:

```

1      min  mean  max
2 1 0.17 0.173 0.176

```

Other Helper Tools

pbdMPI Also contains useful tools for Manager/Worker and task parallelism codes:

- **Task Subsetting:** Distributing a list of jobs/tasks
`get.jid(n)`
- ***ply:** Functions in the *ply family.
`pbdApply(X, MARGIN, FUN, ...)` — analogue of `apply()`
`pbdLapply(X, FUN, ...)` — analogue of `lapply()`
`pbdSapply(X, FUN, ...)` — analogue of `sapply()`

Quick Comments for Using pbdMPI

- 1 Start by loading the package:

```
1 library(pbdMPI, quiet = TRUE)
```

- 2 Always initialize before starting and finalize when finished:

```
1 init()  
2  
3 # ...  
4  
5 finalize()
```