Programming with Big Data in R

George Ostrouchov

February 17, 2014









Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

The **pbdR** Core Team

Wei-Chen Chen¹ George Ostrouchov² Pragneshkumar Patel³ Drew Schmidt³



Support

This work used resources of National Institute for Computational Sciences at the University of Tennessee, Knoxville, which is supported by the Office of Cyberinfrastructure of the U.S. National Science Foundation under Award No. ARRA-NSF-OCI-0906324 for NICS-RDAV center. This work also used resources of the Oak Ridge Leadership Computing Facility at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-000R22725.



¹Department of Ecology and Evolutionary Biology University of Tennessee, Knoxville TN, USA

²Computer Science and Mathematics Division Oak Ridge National Laboratory, Oak Ridge TN, USA

³ National Institute for Computational Sciences University of Tennessee, Knoxville TN, USA

Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

About This Presentation

Speaking Serial R with a Parallel Accent

The content of this presentation is based in part on the **pbdDEMO** vignette *Speaking Serial R with a Parallel Accent*

http://goo.gl/HZkRt

It contains more examples, and sometimes added detail.



About This Presentation

Installation Instructions

Installation instructions for setting up a pbdR environment are available:

This includes instructions for installing R, MPI, and pbdR.



Contents

- Introduction
- 2 pbdR
- 3 Introduction to pbdMPI
- 4 The Generalized Block Distribution
- Basic Statistics Examples
- 6 Introduction to pbdDMAT and the DMAT Structure
- Examples Using pbdDMAT
- 8 Profiling
- Wrapup



Contents

- Introduction
 - A Concise Introduction to Parallelism
 - Quick Overview of Parallel Hardware



What is Parallelism?

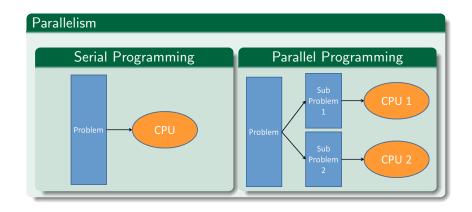
- Doing more than one thing at a time.
- The simultaneous use of multiple compute resources to solve a computational problem.



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○●○000000000
 ○○000
 ○○000
 ○○000
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00
 ○○00

A Concise Introduction to Parallelism





 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 0.000000000
 00000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000</t

A Concise Introduction to Parallelism

Parallelism Serial Programming Parallel Programming make_lunch mpirun -np 2 make_lunch_par Get resources make_lunch_par make lunch par Get resources Get resources Work Work Work Work combine Return Return

A Concise Introduction to Parallelism

Kinds of Parallelism

- Data Parallelism: Data is distributed
- Task Parallelism: Tasks are distributed

(This is a gross oversimplification)



pbdR Paradigms: Data Parallelism

Data parallelism:

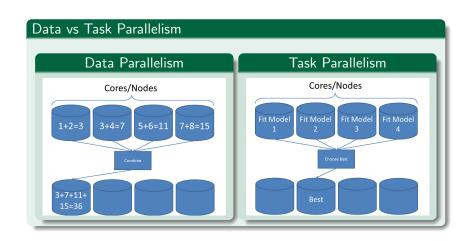
- No one processor/node owns all the data.
- Processors own local pieces of a (conceptually) larger, global object

Task parallelism:

• (Usually) Applying different tasks to the same data.



A Concise Introduction to Parallelism





Parallel Programming Vocabulary: Difficulty in Parallelism

- Implicit parallelism: Parallel details hidden from user
- 2 Explicit parallelism: Some assembly required...
- **3** Embarrassingly Parallel: Also called loosely coupled. Obvious how to make parallel; lots of independence in computations.
- Tightly Coupled: Opposite of embarrassingly parallel; lots of dependence in computations.



Speedup

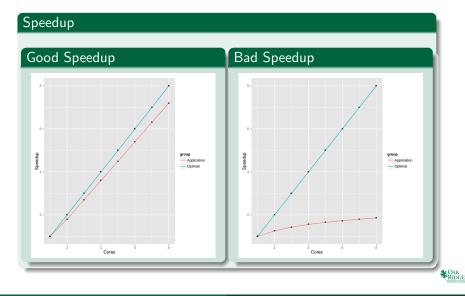
- Wallclock Time: Time of the clock on the wall from start to finish
- Speedup: unitless measure of improvement; more is better.

$$S_{n_1,n_2} = \frac{\text{Run time for } n_1 \text{ cores}}{\text{Run time for } n_2 \text{ cores}}$$

- n_1 is often taken to be 1
- In this case, comparing parallel algorithm to serial algorithm



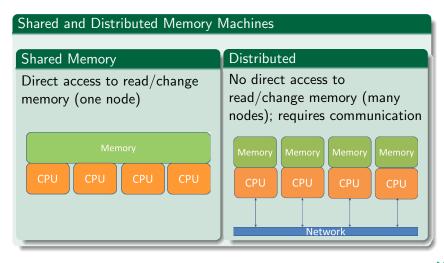
A Concise Introduction to Parallelism



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 00000
 <t

A Concise Introduction to Parallelism





Introduction pbdR IPMbda Stats eg's DMAT pbdDMAT eg's Profiling Wrapup 0000000000

A Concise Introduction to Parallelism

Shared and Distributed Memory Machines

Shared Memory Machines

Thousands of cores



Nautilus. University of Tennessee 1024 cores 4 TB RAM

Distributed Memory Machines

Hundreds of thousands of cores



112.896 cores 147 TB RAM



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○

Quick Overview of Parallel Hardware

Three Basic Flavors of Hardware

| Distributed Memory | Interconnection Network | PROC | PR

Shared Memory



Co-Processor



GPU: Graphical Processing Unit

MIC: Many Integrated Core

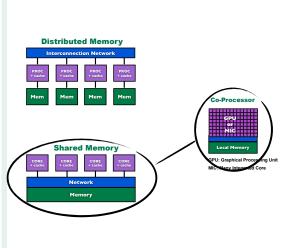


 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○

Quick Overview of Parallel Hardware

Your Laptop or Desktop



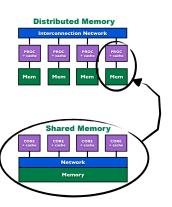


 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 <td

Quick Overview of Parallel Hardware

A Server or Cluster



Co-Processor



GPU: Graphical Processing Unit

MIC: Many Integrated Core

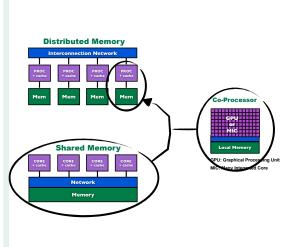


 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 <td

Quick Overview of Parallel Hardware

Server to Supercomputer





 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○

 ○○○○○○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○

Quick Overview of Parallel Hardware

Knowing the Right Words cluster Distributing" Interconnection Network Mem Mem Mem n-Processor Multicore GPU or Manycore "Offloading" **Shared Memory** Network "Multithreading"

OAK RIDGE

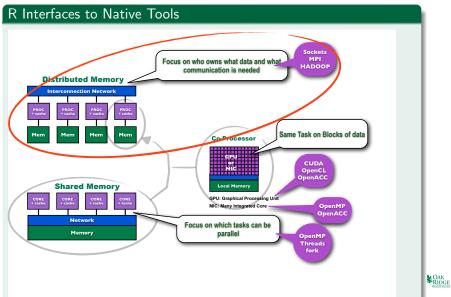
 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 <td

Quick Overview of Parallel Hardware

"Native" Programming Models and Tools Sockets Focus on who owns what data and wha MPI HADOOP communication is needed **Distributed Memory** Same Task on Blocks of data Co-Processor CUDA OpenCL OpenACC **Shared Memory** GPU: Graphical Processing Un MIC: Many Integrated Core OpenMP OpenACC Focus on which tasks can be parallel OpenMP Threads fork OAK RIDGE

Quick Overview of Parallel Hardware



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○●○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○

Quick Overview of Parallel Hardware

30+ Years of Parallel Computing Research Sockets MPI Focus on who owns what data and wha HADOOP communication is needed **Distributed Memory** Interconnection Network Same Task on Blocks of data Mem Mem Mem Mem Co-Processor CUDA MIC OpenCL OpenACC Shared Memory **Local Memory** GPU: Graphical Processing Unit MIC: Many Integrated Core OpenMP OpenACC Network Focus on which tasks can be Memory parallel Open P nreads fork OAK RIDGE

 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○

Quick Overview of Parallel Hardware

Last 10 years of Advances Sockets MPI Focus on who owns what data and wha HADOOP communication is needed **Distributed Memory** Interconnection Network Same Task on Blocks of data Mem Mem Mem Mem Co-Processor CUDA MIC OpenCL OpenACC Shared Memory **Local Memory** GPU: Graphical Processing Unit MIC: Many Integrated Core OpenMP OpenACC Network Focus on which tasks can be Memory parallel Open P nreads fork OAK RIDGE

 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○

Quick Overview of Parallel Hardware

Putting It All Together Challenge Sockets snow Rmpi R Focus on who owns what data and what HADOOP pbdMPI communication is needed **Distributed Memory** RHIPE Same Task on Blocks of data Co-Processor CUDA OpenCL OpenACC **Shared Memory** .c .Call GPU: Graphical Processing Unit MIC: Many Integrated Core -Rcpp OpenMP OpenCL OpenACC inline Focus on which tasks can be ® parallel **OpenMP** Threads fork multicore snow + multicore = parallel (fork) OAK RIDGE R *OAK

 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 <td

Quick Overview of Parallel Hardware

pbdR Focus on Data Parallelism snow R pbdMPI LAPACK Focus on who owns what data and Sockets BLAS MPI what communication is needed Hadoop **Distributed Memory** RHIPE ScaLAPACK **PBLAS BLACS** PETSc Same Task on Co-Processor Trilinos Blocks of data CUDA penCL CUBLAS nACC **Shared Memory** .Call MAGMA Rcpp GPU: Graphical Processin **OpenMP** OpenCL MIC: Many Integrated Core penACC inline MKL ACML Focus on which LibSci OpenMP tasks can be parallel Threads **DPLASMA** fork multicore **PLASMA** snow + multicore = parallel (fork) (OAK RIDGE R *OAK

Contents

- 2 pbdR
 - The pbdR Project
 - Using pbdR





The pbdR Project

Programming with Big Data in R (pbdR)

Striving for Productivity, Portability, Performance



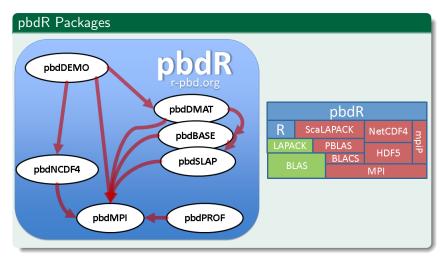
- Free^a R packages.
- Bridging high-performance compiled code with high-productivity of R
- Scalable, big data analytics.
- Offers implicit and explicit parallelism.
- Methods have syntax identical to R.

^aMPL, BSD, and GPL licensed



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

The pbdR Project







The pbdR Project

Simple, Intuitive MPI Operations with pbdMPI

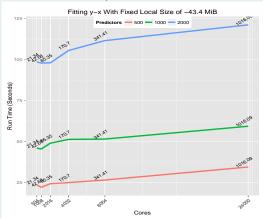
Placeholder



The pbdR Project

Distributed Matrices and Statistics with **pbdDMAT**

Least Squares Benchmark



x <- ddmatrix("rnorm", nrow=m, ncol=n)
y <- ddmatrix("rnorm", nrow=m, ncol=1)
mdl <- lm.fit(x=x, y=y)



http://r-pbd.org/tutorial George Ostrouchov pbdR 26/109

The pbdR Project

Profiling with **pbdPROF**

1. Rebuild **pbdR** packages

```
R CMD INSTALL
    pbdMPI_0.2-1.tar.gz \
    --configure-args= \
    "--enable-pbdPROF"
```

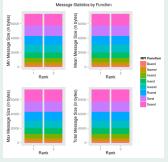
2. Run code

```
mpirun -np 64 Rscript
  my_script.R
```

3. Analyze results

```
library(pbdPROF)
prof <- read.prof(
     "profiler_output.mpiP")
plot(prof)</pre>
```

Publication-quality graphs







pbdR Paradigms

pbdR programs are R programs!

Differences:

- Batch execution (non-interactive).
- Parallel code utilizes Single Program/Multiple Data (SPMD) style
- Emphasizes data parallelism.



Jaten Execution

• Running a serial R program in batch:

or

• Running a parallel (with MPI) R program in batch:

```
mpirun -np 2 Rscript my_par_script.r
```



Single Program/Multiple Data (SPMD)

- SPMD is a programming paradigm.
- Not to be confused with SIMD.

Paradigms

Programming models

e.g. Procedural, OOP, Functional, SPMD, . . .

SIMD

Hardware instructions

e.g. MMX, SSE, ...

SPMD is arguably the simplest extension of serial programming.



- Only one program is written, executed in batch on all processors.
- Different processors are autonomous; there is no manager.
- The dominant programming model for large machines.



Contents

- Introduction to pbdMPI
 - Managing a Communicator
 - Reduce, Gather, Broadcast, and Barrier
 - Other pbdMPI Tools



Message Passing Interface (MPI)

- MPI: Standard for managing communications (data and instructions) between different nodes/computers.
- Implementations: OpenMPI, MPICH2, Cray MPT, . . .
- Enables parallelism (via communication) on distributed machines.
- Communicator: manages communications between processors.



MPI Operations (1 of 2)

 Managing a Communicator: Create and destroy communicators.

```
init() — initialize communicator
finalize() — shut down communicator(s)
```

 Rank query: determine the processor's position in the communicator.

```
comm.rank() — "who am I?"
comm.size() — "how many of us are there?"
```

• **Printing**: Printing output from various ranks.

```
comm.print(x)
comm.cat(x)
```

WARNING: only use these functions on *results*, never on yet-to-be-computed things.



Managing a Communicator

Quick Example 1

Rank Query: 1_rank.r

```
library(pbdMPI, quietly = TRUE)
  init()
3
  my.rank <- comm.rank()
  comm.print(my.rank, all.rank=TRUE)
6
  finalize()
```

Execute this script via:

mpirun -np 2 Rscript 1_rank.r

Sample Output:

```
COMM \cdot RANK = O
  [1] 0
2
  COMM.RANK = 1
  [1]
```



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○

Managing a Communicator

Quick Example 2

Hello World: 2_hello.r

```
library(pbdMPI, quietly=TRUE)
init()

comm.print("Hello, world")

comm.print("Hello again", all.rank=TRUE, quietly=TRUE)

finalize()
```

Execute this script via:

mpirun -np 2 Rscript 2_hello.r

```
COMM.RANK = 0
2 [1] "Hello, world"
3 [1] "Hello again"
```

"Hello again"

Sample Output:



[1]

Reduce, Gather, Broadcast, and Barrier

MPI Operations

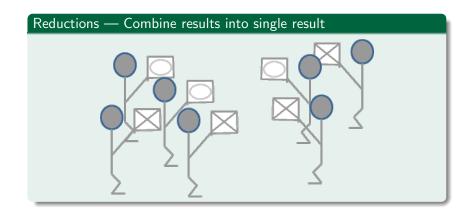
- Reduce
- Gather
- Broadcast
- Barrier



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

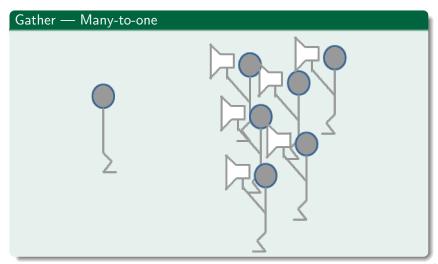
 00000000000
 0000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 <td

Reduce, Gather, Broadcast, and Barrier





Reduce, Gather, Broadcast, and Barrier

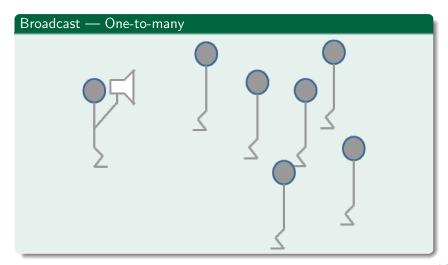




 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○

Reduce, Gather, Broadcast, and Barrier





Reduce, Gather, Broadcast, and Barrier

Barrier — Synchronization Barrier **Barrier** Barrier



MPI Operations (2 of 2)

- Reduction: each processor has a number x; add all of them up, find the largest/smallest,
 reduce(x, op='sum') reduce to one allreduce(x, op='sum') reduce to all
- Gather: each processor has a number; create a new object on some processor containing all of those numbers.
 gather(x) — gather to one
 allgather(x) — gather to all
- Broadcast: one processor has a number x that every other processor should also have.
 bcast(x)
- Barrier: "computation wall"; no processor can proceed until all processors can proceed.
 barrier()



Reduce, Gather, Broadcast, and Barrier

Quick Example 3

```
Reduce and Gather: 3_gt.r
```

```
library(pbdMPI, quietly=TRUE)
  init()
  comm.set.seed(diff=TRUE)
  n <- sample(1:10, size=1)</pre>
  gt <- gather(n)
  comm.print(unlist(gt))
10
  sm <- allreduce(n, op='sum')</pre>
  comm.print(sm, all.rank=T)
13
  finalize()
```

Execute this script via:

Sample Output:

```
COMM.RANK = O
mpirun -np 2 Rscript 3 gt.r
                                   2 [1] 2 8
                                     COMM.RANK = O
                                   4 [1] 10
                                     COMM.RANK = 1
                                   6 [1] 10
```



Reduce, Gather, Broadcast, and Barrier

Quick Example 4

Broadcast: 4_bcast.r

```
library(pbdMPI, quietly=T)
  init()
3
  if (comm.rank() == 0) {
5
    x <- matrix(1:4, nrow=2)
  } else {
    x <- NULL
8
  }
9
10
    <- bcast(x, rank.source=0)
11
  comm.print(y, rank=1)
13
  finalize()
```

Execute this script via:

```
mpirun -np 2 Rscript 4_bcast.r
```

Sample Output:

```
1 COMM.RANK = 1
2 [,1] [,2]
3 [1,] 1 3
4 [2,] 2 4
```



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

Other pbdMPI Tools

MPI Package Controls

•00000

The .SPMD.CT object allows for setting different package options with **pbdMPI**. See the entry *SPMD Control* of the **pbdMPI** manual for information about the .SPMD.CT object:

http://cran.r-project.org/web/packages/pbdMPI/pbdMPI.pdf



Other pbdMPI Tools

Quick Example 5

```
library(pbdMPI, quiet = TRUE)
init()

3
    .SPMD.CT$msg.barrier <- TRUE
5    .SPMD.CT$print.quiet <- TRUE
6
6
6
for (rank in 1:comm.size()-1){
    if (comm.rank() == rank){
        cat(paste("Hello", rank+1, "of", comm.size(), "\n"))
    }
    barrier()
}

12
}
13
14 comm.cat("\n")</pre>
```

Barrier: 5_barrier.r

```
Execute this script via:
```

comm.cat(paste("Hello", comm.rank()+1, "of",
 comm.size(), "\n"), all.rank=TRUE)

Sample Output:

```
1 mpirun -np 2 Rscript 5_barrier.r 1 Hello 1 of 2 Hello 2 of 2
```



finalize()

15

17

Random Seeds

000000

pbdMPI offers a simple interface for managing random seeds:

- comm.set.seed(diff=TRUE) Independent streams via the rlecuyer package.
- comm.set.seed(seed=1234, diff=FALSE) All processors use the same seed seed=1234
- comm.set.seed(diff=FALSE) All processors use the same seed, determined by processor 0 (using the system clock and PID of processor 0).



Other pbdMPI Tools

Quick Example 6

```
Timing: 6_timer.r
```

```
library(pbdMPI, quiet=TRUE)
  init()
  comm.set.seed(diff=T)
  test <- function(timed)
7
    ltime <- system.time(timed)[3]
8
9
10
    mintime <- allreduce(ltime, op='min')
    maxtime <- allreduce(ltime, op='max')
11
    meantime <- allreduce(ltime, op='sum')/comm.size()
12
13
14
    return (data.frame (min=mintime, mean=meantime,
        max=maxtime))
15
16
  times <- test(rnorm(1e6)) # ~7.6MiB of data
  comm.print(times)
19
  finalize()
```

Execute this script via:

```
mpirun -np 2 Rscript 6_timer.r
```

Sample Output:

min mean max 2 1 0.17 0.173 0.176



Other Helper Tools

000000

pbdMPI Also contains useful tools for Manager/Worker and task parallelism codes:

- Task Subsetting: Distributing a list of jobs/tasks get.jid(n)
- *ply: Functions in the *ply family.

 pbdApply(X, MARGIN, FUN, ...) analogue of apply()

 pbdLapply(X, FUN, ...) analogue of lapply()

 pbdSapply(X, FUN, ...) analogue of sapply()



Other pbdMPI Tools

Quick Comments for Using pbdMPI

000000

Start by loading the package:

```
1 library(pbdMPI, quiet = TRUE)
```

Always initialize before starting and finalize when finished:

```
1 init()
2
3 # ...
4
5 finalize()
```



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 0000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 <td

Contents

- 4 The Generalized Block Distribution
 - The GBD Data Structure
 - Example GBD Distributions



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

The GBD Data Structure

Distributing Data

Problem: How to distribute the data

$$x = \begin{bmatrix} x_{1,1} & x_{1,2} & x_{1,3} \\ x_{2,1} & x_{2,2} & x_{2,3} \\ x_{3,1} & x_{3,2} & x_{3,3} \\ x_{4,1} & x_{4,2} & x_{4,3} \\ x_{5,1} & x_{5,2} & x_{5,3} \\ x_{6,1} & x_{6,2} & x_{6,3} \\ x_{7,1} & x_{7,2} & x_{7,3} \\ x_{8,1} & x_{8,2} & x_{8,3} \\ x_{9,1} & x_{9,2} & x_{9,3} \\ x_{10,1} & x_{10,2} & x_{10,3} \end{bmatrix}_{1}$$

7



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

The GBD Data Structure

Distributing a Matrix Across 4 Processors: Block Distribution Data **Processors** 0 $x_{1,1}$ $X_{1,3}$ $X_{1,2}$ $x_{2,2}$ $x_{2,3}$ $X_{2,1}$ X3.1X3,2 X3.3 X4.1X4.2 X4.3 $X_{5.1}$ X5.2 $X_{5.3}$ x = $X_{6,2}$ $X_{6,3}$ $x_{6.1}$ X7.2 $X_{7,3}$ X8.1 X8.2 X8,3 X9.1 X9.2 X9,3



 $X_{10,1}$

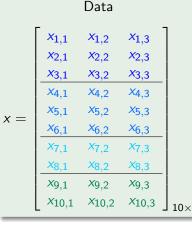
 $X_{10,3}$

*X*_{10,2}

Introduction pbdR pbdMPI **GBD** Stats eg's DMAT pbdDMAT eg's **Profiling** Wrapup 00000

The GBD Data Structure

Distributing a Matrix Across 4 Processors: Local Load Balance



Processors

0



The GBD Data Structure

The GBD Data Structure

Throughout the examples, we will make use of the Generalized Block Distribution, or GBD distributed matrix structure.

 $x_{1.1}$

 $X_{2.1}$

X3.1

X4.1

X5.1

 $x_{6,1}$

X7.1

X8.1

X9.1

X10.1

X1.2

 $X_{2,2}$

X3.2

X4.2

X5.2

 $x_{6,2}$

X7.2

X8.2

X9 2

 $X_{10.2}$

 $x_{1.3}$

 $X_{2.3}$

X3,3

X4.3

X5,3

 $x_{6.3}$

X7.3

X8.3

X9.3

X10.3

- GBD is distributed. No processor owns all the data.
- ② GBD is non-overlapping. Rows uniquely assigned to processors.
- (3) GBD is *row-contiguous*. If a processor owns one element of a row, it owns the entire row.
- 4 GBD is globally row-major, locally column-major.
- GBD is often locally balanced, where each processor owns (almost) the same amount of data. But this is not required.
- The last row of the local storage of a processor is adjacent (by global row) to the first row of the local storage of next processor (by communicator number) that owns data.
- GBD is (relatively) easy to understand, but can lead to bottlenecks if you have many more columns than rows.



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

The GBD Data Structure

Quick Comment for GBD

Local pieces of GBD distributed objects will be given the suffix .gbd to visually help distinguish them from global objects. This suffix carries no semantic meaning.



Example GBD Distributions

Understanding GBD: Global Matrix

$$X = \begin{bmatrix} X_{11} & X_{12} & X_{13} & X_{14} & X_{15} & X_{16} & X_{17} & X_{18} & X_{19} \\ X_{21} & X_{22} & X_{23} & X_{24} & X_{25} & X_{26} & X_{27} & X_{28} & X_{29} \\ X_{31} & X_{32} & X_{33} & X_{34} & X_{35} & X_{36} & X_{37} & X_{38} & X_{39} \\ X_{41} & X_{42} & X_{43} & X_{44} & X_{45} & X_{46} & X_{47} & X_{48} & X_{49} \\ X_{51} & X_{52} & X_{53} & X_{54} & X_{55} & X_{56} & X_{57} & X_{58} & X_{59} \\ X_{61} & X_{62} & X_{63} & X_{64} & X_{65} & X_{66} & X_{67} & X_{68} & X_{69} \\ X_{71} & X_{72} & X_{73} & X_{74} & X_{75} & X_{76} & X_{77} & X_{78} & X_{79} \\ X_{81} & X_{82} & X_{83} & X_{84} & X_{85} & X_{86} & X_{87} & X_{88} & X_{89} \\ X_{91} & X_{92} & X_{93} & X_{94} & X_{95} & X_{96} & X_{97} & X_{98} & X_{99} \end{bmatrix}$$

Processors = 0 1 2 3 4 5



http://r-pbd.org/tutorial

Example GBD Distributions

Understanding GBD: Load Balanced GBD

$$X = \begin{bmatrix} X_{11} & X_{12} & X_{13} & X_{14} & X_{15} & X_{16} & X_{17} & X_{18} & X_{19} \\ X_{21} & X_{22} & X_{23} & X_{24} & X_{25} & X_{26} & X_{27} & X_{28} & X_{29} \\ \hline X_{31} & X_{32} & X_{33} & X_{34} & X_{35} & X_{36} & X_{37} & X_{38} & X_{39} \\ X_{41} & X_{42} & X_{43} & X_{44} & X_{45} & X_{46} & X_{47} & X_{48} & X_{49} \\ \hline X_{51} & X_{52} & X_{53} & X_{54} & X_{55} & X_{56} & X_{57} & X_{58} & X_{59} \\ \hline X_{61} & X_{62} & X_{63} & X_{64} & X_{65} & X_{66} & X_{67} & X_{68} & X_{69} \\ \hline X_{71} & X_{72} & X_{73} & X_{74} & X_{75} & X_{76} & X_{77} & X_{78} & X_{79} \\ \hline X_{81} & X_{82} & X_{83} & X_{84} & X_{85} & X_{86} & X_{87} & X_{88} & X_{89} \\ \hline X_{91} & X_{92} & X_{93} & X_{94} & X_{95} & X_{96} & X_{97} & X_{98} & X_{99} \end{bmatrix}$$

Processors = 0 1 2 3 4 5



Example GBD Distributions

Understanding GBD: Local View

```
X<sub>12</sub>
                 X13
                         X14
                                X<sub>15</sub>
                                        X16
                                                X17
                                                        X<sub>18</sub>
                                                               X19
  x_{21}
         X22
                 X23
                         X24
                                X25
                                        X26
                                                X27
                                                        X28
                                                               X29
  X31
         X32
                 X33
                         X34
                                X35
                                        X36
                                                X37
                                                        X38
                                                               X39
  X41
         X42
                 X43
                         X44
                                X45
                                        X46
                                                X47
                                                        X48
                                                               X49
                         X54
                                X55
                                        X56
                                                X57
                                                        X58
                                                               X59
         X62
                         X<sub>64</sub>
  X_{61}
                 X63
                                X<sub>65</sub>
                                        X66
                                                X67
                                                        X68
                                                               X69
X<sub>71</sub>
          X72
                                                                X79
                 X73
                         X74
                                 X75
                                         X76
                                                X77
                                                        X78
  X81
          X82
                 X83
                         X84
                                 X85
                                         X86
                                                X87
                                                        X88
  X91
          X92
                 X93
                         X94
                                 X95
                                         X96
                                                X97
                                                        X98
```

Processors = 0 1 2 3 4 5



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 0000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 <td

Contents

- Basic Statistics Examples
 - pbdMPI Example: Monte Carlo Simulation
 - pbdMPI Example: Sample Covariance
 - pbdMPI Example: Linear Regression

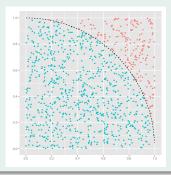


pbdMPI Example: Monte Carlo Simulation

Example 1: Monte Carlo Simulation

Sample *N* uniform observations (x_i, y_i) in the unit square $[0, 1] \times [0, 1]$. Then

$$\pi pprox 4\left(rac{\#\ \textit{Inside Circle}}{\#\ \textit{Total}}
ight) = 4\left(rac{\#\ \mathsf{Blue}}{\#\ \mathsf{Blue} + \#\ \mathsf{Red}}
ight)$$





58/109

pbdMPI Example: Monte Carlo Simulation

Example 1: Monte Carlo Simulation GBD Algorithm

- Let n be big-ish; we'll take n = 50,000.
- **②** Generate an $n \times 2$ matrix x of standard uniform observations.
- **3** Count the number of rows satisfying $x^2 + y^2 \le 1$
- Ask everyone else what their answer is; sum it all up.
- \odot Take this new answer, multiply by 4 and divide by n
- 1 If my rank is 0, print the result.



pbdMPI Example: Monte Carlo Simulation

Example 1: Monte Carlo Simulation Code

Serial Code

```
1 N <- 50000
2 X <- matrix(runif(N * 2), ncol=2)
3 r <- sum(rowSums(X^2) <= 1)
4 PI <- 4*r/N
5 print(PI)</pre>
```

Parallel Code

```
library(pbdMPI, quiet = TRUE)
init()
comm.set.seed(diff=TRUE)

N.gbd <- 50000 / comm.size()
X.gbd <- matrix(runif(N.gbd * 2), ncol = 2)
r.gbd <- sum(rowSums(X.gbd^2) <= 1)
r <- allreduce(r.gbd)
PI <- 4*r/(N.gbd * comm.size())
comm.print(PI)
in finalize()</pre>
```



pbdMPI Example: Monte Carlo Simulation

Note

For the remainder, we will exclude loading, init, and finalize calls.



pbdMPI Example: Sample Covariance

Example 2: Sample Covariance

$$cov(x_{n \times p}) = \frac{1}{n-1} \sum_{i=1}^{n} (x_i - \mu_x) (x_i - \mu_x)^T$$



Example 2: Sample Covariance GBD Algorithm

- \bullet Determine the total number of rows N.
- 2 Compute the vector of column means of the full matrix.
- Subtract each column's mean from that column's entries in each local matrix.
- Ompute the crossproduct locally and reduce.
- **1** Divide by N-1.



pbdMPI Example: Sample Covariance

Example 2: Sample Covariance Code

Serial Code

```
1  N <- nrow(X)
2  mu <- colSums(X) / N
3
4  X <- sweep(X, STATS=mu, MARGIN=2)
5  Cov.X <- crossprod(X) / (N-1)
6
7  print(Cov.X)</pre>
```

Parallel Code

```
1 N <- allreduce(nrow(X.gbd), op="sum")
2 mu <- allreduce(colSums(X.gbd) / N, op="sum")
3 
4 X.gbd <- sweep(X.gbd, STATS=mu, MARGIN=2)
5 Cov.X <- allreduce(crossprod(X.gbd), op="sum") / (N-1)
6 
7 comm.print(Cov.X)</pre>
```



•00

pbdMPI Example: Linear Regression

Example 3: Linear Regression

Find β such that

$$\mathsf{y} = \mathsf{X} eta + \epsilon$$

When **X** is full rank,

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$



pbdMPI Example: Linear Regression

Example 3: Linear Regression GBD Algorithm

- Locally, compute $tx = x^T$
- ② Locally, compute A = tx * x. Query every other processor for this result and sum up all the results.
- 3 Locally, compute B = tx * y. Query every other processor for this result and sum up all the results.
- **1** Locally, compute $A^{-1} * B$



pbdMPI Example: Linear Regression

Example 3: Linear Regression Code

Serial Code

```
1 tX <- t(X)
2 A <- tX %*% X
3 B <- tX %*% y
4 ols <- solve(A) %*% B
```

Parallel Code

```
tX.gbd <- t(X.gbd)
tX.gbd <- t(X.gbd)
tX.gbd <- t(X.gbd %*% X.gbd, op = "sum")

B <- allreduce(tX.gbd %*% y.gbd, op = "sum")

ols <- solve(A) %*% B</pre>
```



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000

Contents

- 6 Introduction to pbdDMAT and the DMAT Structure
 - Introduction to Distributed Matrices
 - DMAT Distributions
 - pbdDMAT



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

Introduction to Distributed Matrices

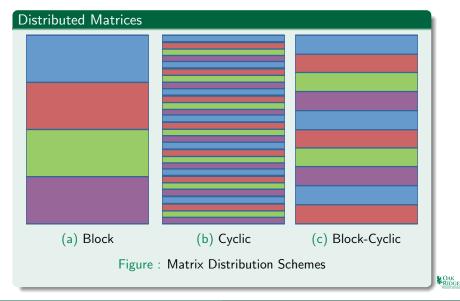
Distributed Matrices

Most problems in data science are matrix algebra problems, so:

Distributed matrices ⇒ Handle Bigger data



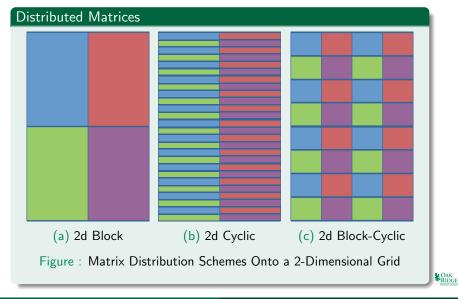
Introduction to Distributed Matrices



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 00000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 <

Introduction to Distributed Matrices



Introduction to Distributed Matrices

Processor Grid Shapes

$$\begin{bmatrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{bmatrix}^{T} \qquad \begin{bmatrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 & 2 \\ 3 & 4 & 5 \end{bmatrix} \qquad \begin{bmatrix} 0 & 1 \\ 2 & 3 \\ 4 & 5 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 & 2 \\ 3 & 4 & 5 \end{bmatrix} \qquad \begin{bmatrix} 0 & 1 \\ 2 & 3 \\ 4 & 5 \end{bmatrix}$$
(a) 1×6 (b) 2×3 (c) 3×2 (d) 6×1

Table: Processor Grid Shapes with 6 Processors



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 0000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 <td

DMAT Distributions

Understanding Dmat: Global Matrix X11 X₁₂ X₁₃ X14 X₁₅ X16 X17 X₁₈ X19 X21 X22 X23 X24 X25 X26 X27 X28 X29 X31 X32 X33 X34 X35 X36 X37 X38 X39 X46 X_{41} X42 X43 X44 X45 X47 X48 X49 x = X_{51} X₅₂ X53 X54 X55 *X*56 X57 X58 X59 X_{61} X62 X63 X64 *X*65 X66 X67 X68 *X*69 X72 X79 X71 *X*73 X74 *X*75 *X*76 *X*77 X78 X88 X81 X82 X83 X84 X85 X86 *X*87 *X*89 X91 X92 *X*93 X94 X95 X96 X97 *X*98 *X*99



DMAT Distributions

DMAT: 1-dimensional Row Block

$$x = \begin{bmatrix} x_{11} & x_{12} & x_{13} & x_{14} & x_{15} & x_{16} & x_{17} & x_{18} & x_{19} \\ x_{21} & x_{22} & x_{23} & x_{24} & x_{25} & x_{26} & x_{27} & x_{28} & x_{29} \\ x_{31} & x_{32} & x_{33} & x_{34} & x_{35} & x_{36} & x_{37} & x_{38} & x_{39} \\ \hline x_{41} & x_{42} & x_{43} & x_{44} & x_{45} & x_{46} & x_{47} & x_{48} & x_{49} \\ x_{51} & x_{52} & x_{53} & x_{54} & x_{55} & x_{56} & x_{57} & x_{58} & x_{59} \\ \hline x_{61} & x_{62} & x_{63} & x_{64} & x_{65} & x_{66} & x_{67} & x_{68} & x_{69} \\ \hline x_{71} & x_{72} & x_{73} & x_{74} & x_{75} & x_{76} & x_{77} & x_{78} & x_{79} \\ x_{81} & x_{82} & x_{83} & x_{84} & x_{85} & x_{86} & x_{87} & x_{88} & x_{89} \\ x_{91} & x_{92} & x_{93} & x_{94} & x_{95} & x_{96} & x_{97} & x_{98} & x_{99} \end{bmatrix}$$

Processor grid =
$$\begin{vmatrix} 0 \\ 1 \\ 2 \\ 3 \end{vmatrix} = \begin{vmatrix} (0,0) \\ (1,0) \\ (2,0) \\ (3,0) \end{vmatrix}$$



http://r-pbd.org/tutorial

DMAT Distributions

DMAT: 2-dimensional Row Block

$$X = \begin{bmatrix} X_{11} & X_{12} & X_{13} & X_{14} & X_{15} & X_{16} & X_{17} & X_{18} & X_{19} \\ X_{21} & X_{22} & X_{23} & X_{24} & X_{25} & X_{26} & X_{27} & X_{28} & X_{29} \\ X_{31} & X_{32} & X_{33} & X_{34} & X_{35} & X_{36} & X_{37} & X_{38} & X_{39} \\ X_{41} & X_{42} & X_{43} & X_{44} & X_{45} & X_{46} & X_{47} & X_{48} & X_{49} \\ X_{51} & X_{52} & X_{53} & X_{54} & X_{55} & X_{56} & X_{57} & X_{58} & X_{59} \\ \hline X_{61} & X_{62} & X_{63} & X_{64} & X_{65} & X_{66} & X_{67} & X_{68} & X_{69} \\ X_{71} & X_{72} & X_{73} & X_{74} & X_{75} & X_{76} & X_{77} & X_{78} & X_{79} \\ X_{81} & X_{82} & X_{83} & X_{84} & X_{85} & X_{86} & X_{87} & X_{88} & X_{89} \\ X_{91} & X_{92} & X_{93} & X_{94} & X_{95} & X_{96} & X_{97} & X_{98} & X_{99} \end{bmatrix}$$

Processor grid =
$$\begin{vmatrix} 0 & 1 \\ 2 & 3 \end{vmatrix} = \begin{vmatrix} (0,0) & (0,1) \\ (1,0) & (1,1) \end{vmatrix}$$



Introduction pbdR pbdMPI Stats eg's DMAT pbdDMAT eg's Profiling Wrapup 000000

DMAT Distributions

DMAT: 1-dimensional Row Cyclic

$$X = \begin{bmatrix} x_{11} & x_{12} & x_{13} & x_{14} & x_{15} & x_{16} & x_{17} & x_{18} & x_{19} \\ \hline x_{21} & x_{22} & x_{23} & x_{24} & x_{25} & x_{26} & x_{27} & x_{28} & x_{29} \\ \hline x_{31} & x_{32} & x_{33} & x_{34} & x_{35} & x_{36} & x_{37} & x_{38} & x_{39} \\ \hline x_{41} & x_{42} & x_{43} & x_{44} & x_{45} & x_{46} & x_{47} & x_{48} & x_{49} \\ \hline x_{51} & x_{52} & x_{53} & x_{54} & x_{55} & x_{56} & x_{57} & x_{58} & x_{59} \\ \hline x_{61} & x_{62} & x_{63} & x_{64} & x_{65} & x_{66} & x_{67} & x_{68} & x_{69} \\ \hline x_{71} & x_{72} & x_{73} & x_{74} & x_{75} & x_{76} & x_{77} & x_{78} & x_{79} \\ \hline x_{81} & x_{82} & x_{83} & x_{84} & x_{85} & x_{86} & x_{87} & x_{88} & x_{89} \\ \hline x_{91} & x_{92} & x_{93} & x_{94} & x_{95} & x_{96} & x_{97} & x_{98} & x_{99} \end{bmatrix}$$

Processor grid =
$$\begin{vmatrix} 0 \\ 1 \\ 2 \\ 3 \end{vmatrix} = \begin{vmatrix} (0,0) \\ (1,0) \\ (2,0) \\ (3,0) \end{vmatrix}$$



75/109 pbdR

 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 0000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 <td

DMAT Distributions

DMAT: 2-dimensional Row Cyclic

$$X = \begin{bmatrix} x_{11} & x_{12} & x_{13} & x_{14} & x_{15} & x_{16} & x_{17} & x_{18} & x_{19} \\ \hline x_{21} & x_{22} & x_{23} & x_{24} & x_{25} & x_{26} & x_{27} & x_{28} & x_{29} \\ \hline x_{31} & x_{32} & x_{33} & x_{34} & x_{35} & x_{36} & x_{37} & x_{38} & x_{39} \\ \hline x_{41} & x_{42} & x_{43} & x_{44} & x_{45} & x_{46} & x_{47} & x_{48} & x_{49} \\ \hline x_{51} & x_{52} & x_{53} & x_{54} & x_{55} & x_{56} & x_{57} & x_{58} & x_{59} \\ \hline x_{61} & x_{62} & x_{63} & x_{64} & x_{65} & x_{66} & x_{67} & x_{68} & x_{69} \\ \hline x_{71} & x_{72} & x_{73} & x_{74} & x_{75} & x_{76} & x_{77} & x_{78} & x_{79} \\ \hline x_{81} & x_{82} & x_{83} & x_{84} & x_{85} & x_{86} & x_{87} & x_{88} & x_{89} \\ \hline x_{91} & x_{92} & x_{93} & x_{94} & x_{95} & x_{96} & x_{97} & x_{98} & x_{99} \end{bmatrix}$$

Processor grid =
$$\begin{vmatrix} 0 & 1 \\ 2 & 3 \end{vmatrix} = \begin{vmatrix} (0,0) & (0,1) \\ (1,0) & (1,1) \end{vmatrix}$$



http://r-pbd.org/tutorial

DMAT Distributions

DMAT: 2-dimensional Block-Cyclic

$$x = \begin{bmatrix} x_{11} & x_{12} & x_{13} & x_{14} & x_{15} & x_{16} & x_{17} & x_{18} & x_{19} \\ x_{21} & x_{22} & x_{23} & x_{24} & x_{25} & x_{26} & x_{27} & x_{28} & x_{29} \\ x_{31} & x_{32} & x_{33} & x_{34} & x_{35} & x_{36} & x_{37} & x_{38} & x_{39} \\ x_{41} & x_{42} & x_{43} & x_{44} & x_{45} & x_{46} & x_{47} & x_{48} & x_{49} \\ x_{51} & x_{52} & x_{53} & x_{54} & x_{55} & x_{56} & x_{57} & x_{58} & x_{59} \\ x_{61} & x_{62} & x_{63} & x_{64} & x_{65} & x_{66} & x_{67} & x_{68} & x_{69} \\ x_{71} & x_{72} & x_{73} & x_{74} & x_{75} & x_{76} & x_{77} & x_{78} & x_{79} \\ x_{81} & x_{82} & x_{83} & x_{84} & x_{85} & x_{86} & x_{87} & x_{88} & x_{89} \\ x_{91} & x_{92} & x_{93} & x_{94} & x_{95} & x_{96} & x_{97} & x_{98} & x_{99} \end{bmatrix}$$

Processor grid =
$$\begin{vmatrix} 0 & 1 \\ 2 & 3 \end{vmatrix} = \begin{vmatrix} (0,0) & (0,1) \\ (1,0) & (1,1) \end{vmatrix}$$



http://r-pbd.org/tutorial

Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

•0000000

pbdDMAT

The DMAT Data Structure

The more complicated the processor grid, the more complicated the distribution.



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

0000000

pbdDMAT

DMAT: 2-dimensional Block-Cyclic with 6 Processors

$$X = \begin{bmatrix} X_{11} & X_{12} & X_{13} & X_{14} & X_{15} & X_{16} & X_{17} & X_{18} & X_{19} \\ X_{21} & X_{22} & X_{23} & X_{24} & X_{25} & X_{26} & X_{27} & X_{28} & X_{29} \\ X_{31} & X_{32} & X_{33} & X_{34} & X_{35} & X_{36} & X_{37} & X_{38} & X_{39} \\ X_{41} & X_{42} & X_{43} & X_{44} & X_{45} & X_{46} & X_{47} & X_{48} & X_{49} \\ X_{51} & X_{52} & X_{53} & X_{54} & X_{55} & X_{56} & X_{57} & X_{58} & X_{59} \\ X_{61} & X_{62} & X_{63} & X_{64} & X_{65} & X_{66} & X_{67} & X_{68} & X_{69} \\ X_{71} & X_{72} & X_{73} & X_{74} & X_{75} & X_{76} & X_{77} & X_{78} & X_{79} \\ X_{81} & X_{82} & X_{83} & X_{84} & X_{85} & X_{86} & X_{87} & X_{88} & X_{89} \\ \hline X_{91} & X_{92} & X_{93} & X_{94} & X_{95} & X_{96} & X_{97} & X_{98} & X_{99} \end{bmatrix}$$

Processor grid =
$$\begin{vmatrix} 0 & 1 & 2 \\ 3 & 4 & 5 \end{vmatrix} = \begin{vmatrix} (0,0) & (0,1) & (0,2) \\ (1,0) & (1,1) & (1,2) \end{vmatrix}$$



http://r-pbd.org/tutorial

Introduction pbdR pbdMPI Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

*X*₁₃

X23

X53

X63

X93

X33

X43

X73

X83

X₁₄

X24

X54

X₆₄

X94

X34

 X_{44}

X74

X84

00000000

X19

X29

X59

X₆₉

*X*99

X39

X49

X79

X89

pbdDMAT

Understanding DMAT: Local View

Processor grid =

$$\begin{vmatrix} 1 & 2 \\ 4 & 5 \end{vmatrix} =$$

$$(0,1)$$
 $(0,1)$ $(1,1)$

X16

 X_{26}

X56

X66

X96

X36

X46

X76

X86

X55

 X_{65}

X95

X35

X₄₅



http://r-pbd.org/tutorial

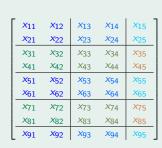
Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

00000000

pbdDMAT

The DMAT Data Structure

- ① DMAT is distributed. No one processor owns all of the matrix.
- ② DMAT is non-overlapping. Any piece owned by one processor is owned by no other processors.
- ① DMAT can be row-contiguous or not, depending on the processor grid and blocking factor used.
- OMAT is locally column-major and globally, it depends...
- GBD is a generalization of the one-dimensional block DMAT distribution. Otherwise there is no relation.
- DMAT is confusing, but very robust.





Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

pbdDMAT

Pros and Cons of This Data Structure

Pros

 Robust for matrix computations.

Cons

00000000

Confusing layout.

This is why we hide most of the distributed details.

The details are there if you want them (you don't want them).



pbdDMAT

Distributed Matrix Methods

pbdDMAT has over 100 methods with *identical* syntax to R:

- `[`, rbind(), cbind(), ...
- lm.fit(), prcomp(), cov(), ...
- `%*%`, solve(), svd(), norm(), ...
- median(), mean(), rowSums(), ...

Serial Code

1 cov(x)

Parallel Code

cov(x)



Comparing pbdMPI and pbdDMAT

pbdMPI:

- MPI + sugar.
- GBD not the only structure pbdMPI can handle (just a useful convention).

pbdDMAT:

- More of a software package.
- DMAT structure must be used for pbdDMAT.
- If the data is not 2d block-cyclic compatible, DMAT will definitely give the wrong answer.



Quick Comments for Using pbdDMAT

Start by loading the package:

```
library(pbdDMAT, quiet = TRUE)
```

② Always initialize before starting and finalize when finished:

```
init.grid()
3
  finalize()
```

O Distributed DMAT objects will be given the suffix .dmat to visually help distinguish them from global objects. This suffix carries no semantic meaning.



Contents

- Examples Using pbdDMAT
 - Manipulating DMAT Objects
 - Statistics Examples with pbdDMAT
 - RandSVD



Manipulating DMAT Objects

Example 1: DMAT Construction

Generate a global matrix and distribute it

```
library(pbdDMAT, quiet=TRUE)
  init.grid()
3
  # Common global on all processors --> distributed
  x <- matrix(1:100, nrow=10, ncol=10)
  x.dmat <- as.ddmatrix(x)
8 x.dmat
9
  # Global on processor 0 --> distributed
  if (comm.rank() == 0) {
    v <- matrix(1:100, nrow=10, ncol=10)</pre>
12
  } else {
13
    y <- NULL
14
15
16 y.dmat <- as.ddmatrix(y)
17
  y.dmat
18
19
20 finalize()
```

```
mpirun -np 2 Rscript 1-gen.r
```



Introduction

Example 2: DMAT Construction

Generate locally only what is needed

```
library(pbdDMAT, quiet=TRUE)
  init.grid()
3
  zero.dmat <- ddmatrix(0, nrow=100, ncol=100)
  zero.dmat
6
  id.dmat <- diag(1, nrow=100, ncol=100, type="ddmatrix")
  id.dmat
9
  comm.set.seed(diff=T)
10
  rand.dmat <- ddmatrix("rnorm", nrow=100, ncol=100.
      mean=10. sd=100)
12
  rand.dmat
13
14
  finalize()
```

```
mpirun -np 2 Rscript 2_gen.r
```



Generate locally only what is needed

```
library(pbdDMAT, quiet=TRUE)
init.grid()

4      x.dmat <- ddmatrix(1:30, nrow=10)
5      y.dmat <- x.dmat[c(1, 3, 5, 7, 9), -3]
6
7      comm.print(y.dmat)
8      y <- as.matrix(y.dmat)
9      comm.print(y)
10
11      finalize()</pre>
```

```
mpirun —np 2 Rscript 3_extract.r
```



Introduction

Example 4: More DMAT Operations

```
library(pbdDMAT, quiet=TRUE)
  init.grid()
3
  x.dmat <- ddmatrix(1:30, nrow=10)
5
  y.dmat <- x.dmat + 1:7
7
  z.dmat <- scale(x.dmat, center=TRUE, scale=TRUE)
9
  y <- as.matrix(y.dmat)
10
  z <- as.matrix(z.dmat)
11
12
13
  comm.print(y)
  comm.print(z)
14
15
  finalize()
16
```

Execute this script via:



http://r-pbd.org/tutorial George Ostrouchov pbdR 89/109

```
1 % mpirun —np 2 Rscript 4_convert.r 2 %
```



Sample Covariance Serial Code Cov.X <- cov(X) print(Cov.X) Parallel Code Cov.X <- cov(X) print(Cov.X)



Linear Regression

Serial Code

```
tX <- t(X)
A <- tX %*% X
B <- tX %*% y

ols <- solve(A) %*% B

# or
ols <- lm.fit(X, y)</pre>
```

Parallel Code

```
1 tX <- t(X)
2 A <- tX %*% X
3 B <- tX %*% y
4
5 ols <- solve(A) %*% B
6
7 # or
8 ols <- lm.fit(X, y)
```



Example 5: PCA

PCA: pca.r

```
library (pbdDMAT, quiet=T)
    init.grid()
2
3
4
5
6
7
   n < -1e4
   p < -250
   comm.set.seed(diff=T)
8
   x.dmat <- ddmatrix("rnorm", nrow=n, ncol=p, mean=100, sd=25)
    pca <- prcomp(x=x.dmat, retx=TRUE, scale=TRUE)</pre>
10
11
    prop_var <- cumsum(pca$sdev)/sum(pca$sdev)</pre>
12
    i \leftarrow max(min(which(prop_var > 0.9)) - 1, 1)
13
14
   v.dmat \leftarrow pca$x[. 1:i]
15
   comm.cat("\nCols: ", i, "\n", quiet=T)
16
   comm.cat("\%Cols:", i/dim(x.dmat)[2], "\n\n", quiet=T)
17
18
19
    finalize()
```

Execute this script via:

Sample Output:

```
1 mpirun —np 2 Rscript 5_pca.r 1 Cols: 221 %Cols: 0.884
```



 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 ○○○○○○○○○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○
 ○○○○
 ○○○○
 ○○○○
 ○○○○

Statistics Examples with pbdDMAT

Distributed Matrices

pbdDEMO contains many other examples of reading and managing GBD and DMAT data



RandSVD

Randomized SVD1

Prototype for Randomized SVD Given an $m \times n$ matrix A, a target number k of singular vectors, and an exponent q (say, q = 1 or q = 2), this procedure computes an approximate rank-2k factorization $U\Sigma V^*$, where U and V are orthonormal, and Σ is nonnegative and diagonal.

Stage A:

- Generate an $n \times 2k$ Gaussian test matrix Ω .
- 2 Form Y = (AA*)^qAΩ by multiplying alternately with A and A*. 3 Construct a matrix Q whose columns form an orthonormal basis for

the range of Y. Stage B:

- 4 Form $B = Q^*A$.
- Compute an SVD of the small matrix: $B = \tilde{U}\Sigma V^*$.
- 6 Set $U = O\widetilde{U}$.

 $Q = Q_a$.

Note: The computation of Y in step 2 is vulnerable to round-off errors. When high accuracy is required, we must incorporate an orthonormalization step between each application of A and A^* ; see Algorithm 4.4.

```
Algorithm 4.4: Randomized Subspace Iteration
Given an m \times n matrix A and integers \ell and q, this algorithm computes an
m \times \ell orthonormal matrix Q whose range approximates the range of A.
    Draw an n \times \ell standard Gaussian matrix \Omega.
    Form Y_0 = A\Omega and compute its OR factorization Y_0 = Q_0R_0.
    for j = 1, 2, ..., q
         Form \tilde{Y}_i = A^*Q_{i-1} and compute its QR factorization \tilde{Y}_i = \tilde{Q}_i\tilde{R}_i.
         Form Y_i = A\widetilde{Q}_i and compute its QR factorization Y_i = Q_iR_i.
6
    end
```

Serial R

•00

```
randSVD \leftarrow function(A, k, g=3)
2
3
        ## Stage A
        Omega <- matrix(rnorm(n*2*k),
4
5
                   nrow=n. ncol=2*k)
6
        Y <- A %*% Omega
7
        Q \leftarrow ar.Q(ar(Y))
8
         At \leftarrow t(A)
9
         for(i in 1:q)
10
              Y <- At %*% O
11
12
             Q \leftarrow qr.Q(qr(Y))
             Y <- A %*% Q
13
              Q \leftarrow ar.Q(ar(Y))
14
15
16
17
        ## Stage B
        B <- t(Q) %*% A
18
19
        U <- La.svd(B)$u
20
        U <- Q %*% U
21
        U[, 1:k]
22
```

¹Halko N, Martinsson P-G and Tropp J A 2011 Finding structure with randomness: probabilistic algorithms for constructing approximate matrix decompositions SIAM Rev. 53 217-88



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's Profiling Wrapup

RandSVD

Randomized SVD

Serial R

```
randSVD \leftarrow function(A, k, q=3)
 2
 3
        ## Stage A
 4
         Omega <- matrix(rnorm(n*2*k),
                nrow=n. ncol=2*k)
 6
         Y <- A %*% Omega
         Q \leftarrow qr.Q(qr(Y))
8
         At \leftarrow t(A)
9
         for(i in 1:q)
10
11
              Y <- At %*% Q
12
             Q \leftarrow qr.Q(qr(Y))
13
              Y <- A %*% Q
14
              Q \leftarrow qr.Q(qr(Y))
15
16
17
        ## Stage B
18
         B <- t(Q) %*% A
19
         U <- La.svd(B)$u
20
         U <- Q %*% U
21
         U[, 1:k]
22
```

Parallel pbdR

000

```
randSVD \leftarrow function(A, k, q=3)
 3
        ## Stage A
         Omega <- ddmatrix("rnorm",
                nrow=n. ncol=2*k)
         Y <- A %*% Omega
        Q \leftarrow qr.Q(qr(Y))
         At \leftarrow t(A)
         for(i in 1:q)
10
11
              Y <- At %*% Q
12
              Q \leftarrow qr.Q(qr(Y))
13
              Y <- A %*% Q
14
              Q \leftarrow qr.Q(qr(Y))
15
16
17
        ## Stage B
18
         B <- t(Q) %*% A
19
         U <- La.svd(B)$u</p>
20
         U <- Q %*% U
21
         U[, 1:k]
22
```

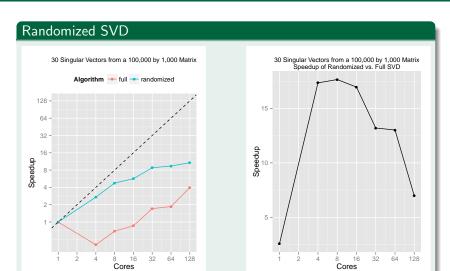


 Introduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 00000000000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000
 0000

000

RandSVD





Contents

- 8 Profiling
 - Profiling
 - Installing pbdPROF
 - Example

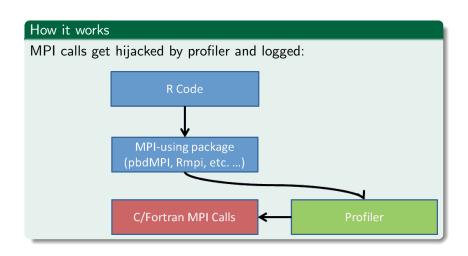


Introduction to pbdPROF

- Successful Google Summer of Code 2013 project.
- Available on the CRAN.
- Enables profiling of MPI-using R scripts.
- pbdR packages officially supported; can work with others. . .
- Also reads, parses, and plots profiler outputs.



Profiling





- Currently supports the profilers fpmpi and mpiP.
- fpmpi is distributed with pbdPROF and installs easily, but offers minimal profiling capabilities.
- mpiP is fully supported also, but harder to install.



Installing pbdPROF

Installing pbdPROF

- Build pbdPROF.
- 2 Rebuild **pbdMPI** (linking with **pbdPROF**).
- 3 Run your analysis as usual.
- 1 Interactively analyze profiler outputs with pbdPROF.



Installing pbdPROF

Build **pbdPROF**

R CMD INSTALL pbdPROF_0.2-1.tar.gz

- The above installs **fpmpi**.
- mpiP can be used if you have a system installation available.
- See package vignette for more details and troubleshooting.



```
R CMD INSTALL pbdMPI_0.2-2.tar.gz
--configure-args="--enable-pbdPROF"
```

- Any package which explicitly links with an MPI library must be rebuilt in this way (pbdMPI, Rmpi, ...).
- Other pbdR packages link with pbdMPI, and so do not need to be rebuilt.
- See **pbdPROF** vignette if something goes wrong.



Example

Example Script

my_svd.r

```
library(pbdMPI, lib.loc="~/R/prof", quietly=TRUE)
library(pbdDMAT, quietly=T)
init.grid()

n <- 1000
x <- ddmatrix("rnorm", n, n)

asdf <- La.svd(x)

finalize()</pre>
```



•00000

Example Script

Run example with 4 ranks:

```
$ mpirun -np 4 Rscript my_svd.r
mpiP:
mpiP: mpiP: mpiP V3.4.0 (Build Feb 14 2014/13:55:39)
mpiP: Direct questions and errors to
        mpip-help@lists.sourceforge.net
mpiP:
Using 2x2 for the default grid size

mpiP:
mpiP: Storing mpiP output in [./R.4.28812.1.mpiP].
mpiP:
```



Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's **Profiling** Wrapup

Example

Read Profiler Data into R

Interactively (or in batch) Read in Profiler Data

```
1 library(pbdPROF)
2 prof.data <- read.prof("R.4.28812.1.mpiP")</pre>
```

Partial Output of Example Data

```
> prof.data
An mpip profiler object:
\lceil \lceil 1 \rceil \rceil
  Task
        AppTime MPITime MPI.
1
           3.68 0.00816 0.22
2
           3.68 0.04890 1.33
           3.69 0.03850 1.04
4
           3.68 0.00904 0.25
          14.70 0.10500 0.71
5
[[2]]
       Lev File.Address Line_Parent_Funct
                                                MPI_Call
           1.400699e + 14
                                    [unknown] Allreduce
           1.400699e+14
                                    [unknown] Allreduce
2
```



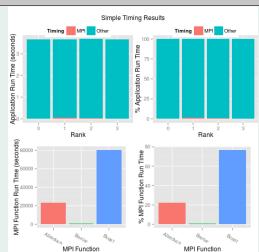
 roduction
 pbdR
 pbdMPI
 GBD
 Stats eg's
 DMAT
 pbdDMAT eg's
 Profiling
 Wrapup

 0000000000
 0000
 0000
 0000
 0000
 0000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 000
 00

Example

Generate plots

plot(prof.data)



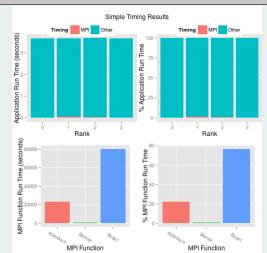


Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's **Profiling** Wrapup

Example

Generate plots

plot(prof.data, plot.type="stats1")



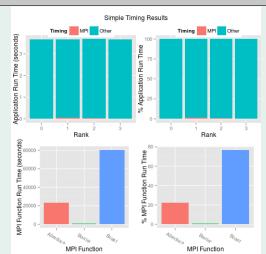


Introduction pbdR pbdMPI GBD Stats eg's DMAT pbdDMAT eg's **Profiling** Wrapup

Example

Generate plots

plot(prof.data, plot.type="messages1")





Contents





The pbdR Project

- Our website: http://r-pbd.org/
- Email us at: RBigData@gmail.com
- Our google group: http://group.r-pbd.org/

Where to begin?

- The pbdDEMO package http://cran.r-project.org/web/packages/pbdDEMO/
- The **pbdDEMO** Vignette: http://goo.gl/HZkRt



Thanks for coming!

Questions?



http://r-pbd.org/

